

Minimum Cost Spanning Trees (MSTs)

Joseph Chuang-Chieh Lin (林莊傑)

Department of Computer Science & Engineering,
National Taiwan Ocean University

Fall 2025



Outline

- 1 Introduction
- 2 Kruskal's algorithm
- 3 Prim's algorithm
- 4 Sollin's/ Borůvka's algorithm
 - Supplementary (Cut Property)

Outline

1 Introduction

2 Kruskal's algorithm

3 Prim's algorithm

4 Sollin's/ Borůvka's algorithm
• Supplementary (Cut Property)

Cost & Minimum-Cost Spanning Tree

- The cost of a spanning tree of a weighted undirected graph is the sum of the weights of the edges in the spanning tree.



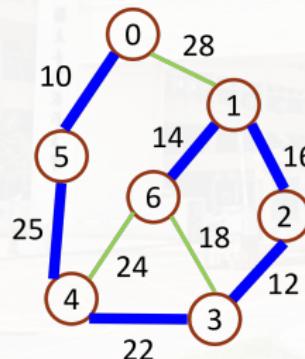
Cost & Minimum-Cost Spanning Tree

- The cost of a spanning tree of a weighted undirected graph is the sum of the weights of the edges in the spanning tree.
- A minimum-cost spanning tree is a spanning tree of least cost.



Cost & Minimum-Cost Spanning Tree

- The cost of a spanning tree of a weighted undirected graph is the sum of the weights of the edges in the spanning tree.
- A minimum-cost spanning tree is a spanning tree of least cost.



$$\text{cost} = 10 + 25 + 22 + 12 + 16 + 14 = 99.$$

$$\text{cost} = 28 + 16 + 12 + 22 + 24 + 25 = 127.$$

Greedy Methods

- We will introduce three different **greedy algorithms** can be used to obtain a minimum cost spanning tree of a connected undirected graph.

Greedy Methods

- We will introduce three different **greedy algorithms** can be used to obtain a minimum cost spanning tree of a connected undirected graph.
 - Kruskal's algorithm
 - Prim's algorithm
 - Sollin's algorithm
- ★ Further reading: Greedy Algorithms & Matroids [[link](#)]

Greedy Method (1/2)

- Construct an optimal solution in stages.
- A feasible solution is one which works within the constraints specified by the problem.
- At each stage, we make a decision that is **the best decision at that time**.
- Typically, the selection of an item at each stage is based on either **a least cost** or **a highest profit** criterion.



Greedy Method (2/2)

- For minimum spanning trees, we use a least cost criterion. Our solution must satisfy the following constraints:

Greedy Method (2/2)

- For minimum spanning trees, we use a least cost criterion. Our solution must satisfy the following constraints:
 - We must use only edges within the graph.



Greedy Method (2/2)

- For minimum spanning trees, we use a least cost criterion. Our solution must satisfy the following constraints:
 - We must use only edges within the graph.
 - Exactly $n - 1$ edges are used.



Greedy Method (2/2)

- For minimum spanning trees, we use a least cost criterion. Our solution must satisfy the following constraints:
 - We must use only edges within the graph.
 - Exactly $n - 1$ edges are used.
 - Never use edges that would produce a cycle.

Outline

1 Introduction

2 Kruskal's algorithm

3 Prim's algorithm

4 Sollin's/ Borůvka's algorithm
• Supplementary (Cut Property)

Kruskal's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.

Kruskal's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.
- Select the edges for inclusion in T in nondecreasing order of their cost.

Kruskal's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.
- Select the edges for inclusion in T in nondecreasing order of their cost.
- An edge is added to T if it does not form a cycle with the edges that are already in T .

Kruskal's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.
- Select the edges for inclusion in T in nondecreasing order of their cost.
- An edge is added to T if it does not form a cycle with the edges that are already in T .
- Exactly $n - 1$ edges will be selected for inclusion in T .



Kruskal's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.
- Select the edges for inclusion in T in nondecreasing order of their cost.
- An edge is added to T if it does not form a cycle with the edges that are already in T .
 - How to do this?
- Exactly $n - 1$ edges will be selected for inclusion in T .



Kruskal's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.
- Select the edges for inclusion in T in nondecreasing order of their cost.
- An edge is added to T if it does not form a cycle with the edges that are already in T .
 - How to do this? \Rightarrow Union-Find Operations!
- Exactly $n - 1$ edges will be selected for inclusion in T .



Kruskal's Algorithm

- Build a minimum cost spanning tree T by adding edges to T one at a time.
- Select the edges for inclusion in T in nondecreasing order of their cost.
- An edge is added to T if it does not form a cycle with the edges that are already in T .
 - How to do this? \Rightarrow Union-Find Operations!
- Exactly $n - 1$ edges will be selected for inclusion in T .
- Time complexity: $O(e \log n)$ (assume using tree-based Union-Find).
 - Sorting the edges: $\approx e \log e < e \log n^2 = O(e \log n)$.
 - At most $2e$ find operations $\approx \log n$ time for each.
 - At most $2n - 1$ union operations $\approx n$ time.



Illustration of Kruskal's Algorithm

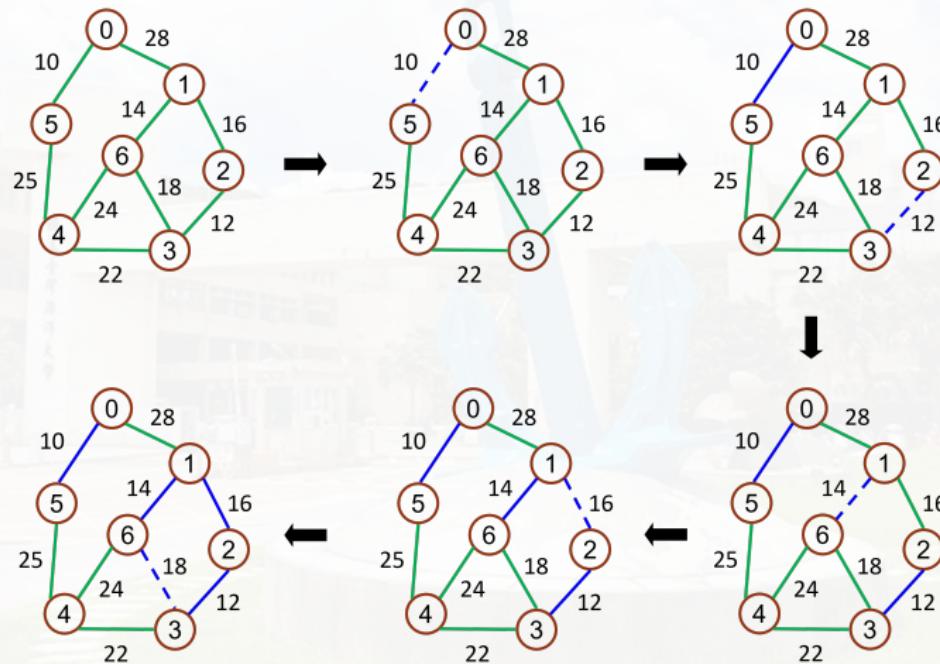
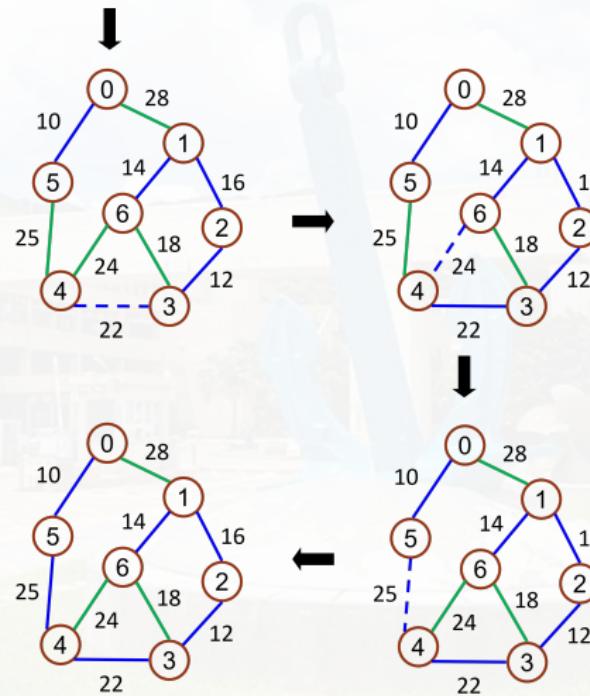


Illustration of Kruskal's Algorithm

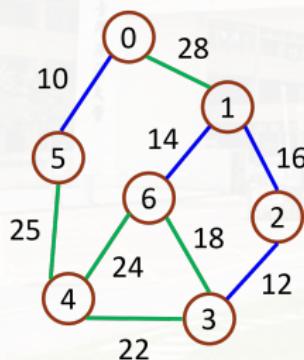


Union-Find Operations

- Union-Find Operations: to determine whether or not adding an edge would cause a cycle.

Union-Find Operations

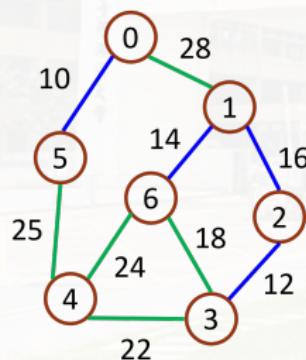
- Union-Find Operations: to determine whether or not adding an edge would cause a cycle.
- For example,



- $\{0, 5\}, \{1, 2, 3, 6\}, \{4\}$: the sets corresponding to existing subtrees.

Union-Find Operations

- Union-Find Operations: to determine whether or not adding an edge would cause a cycle.
- For example,



- $\{0, 5\}, \{1, 2, 3, 6\}, \{4\}$: the sets corresponding to existing subtrees.
- Vertex 3 and 6 are already in the same set \Rightarrow edge $(3, 6)$ is rejected!

The Pseudo-code of Kruskal's Algorithm

```
T = { };
while (T contains fewer than n-1 edges && E is not empty) {
    choose a least cost edge (v,w) from E;
    delete (v,w) from E;
    if ((v,w) does not create a cycle in T)
        add (v,w) to T
    else
        discard (v,w);
}
if (T contains fewer than n-1 edges)
    printf("No spanning tree\n");
```

- How could “No spanning tree” happen?



Outline

- 1 Introduction
- 2 Kruskal's algorithm
- 3 Prim's algorithm
- 4 Sollin's/ Borůvka's algorithm
 - Supplementary (Cut Property)

Prim's Algorithm (1/3)

- Another greedy MST algorithm.



Prim's Algorithm (1/3)

- Another greedy MST algorithm.
- The main difference:

Prim's Algorithm (1/3)

- Another greedy MST algorithm.
- The main difference:
 - The set of selected edges forms a **tree at all times** in Prim's algorithm.
 - The set of selected edges in Kruskal's algorithm forms a *forest at each stage*.

Prim's Algorithm (2/3)

- Prim's algorithm begins with a tree T that contains **one arbitrary single vertex**.

Prim's Algorithm (2/3)

- Prim's algorithm begins with a tree T that contains **one arbitrary single vertex**.
- Next, we add a **least cost edge** (u, v) to T such that $T \cup (u, v)$ is also a tree.

Prim's Algorithm (2/3)

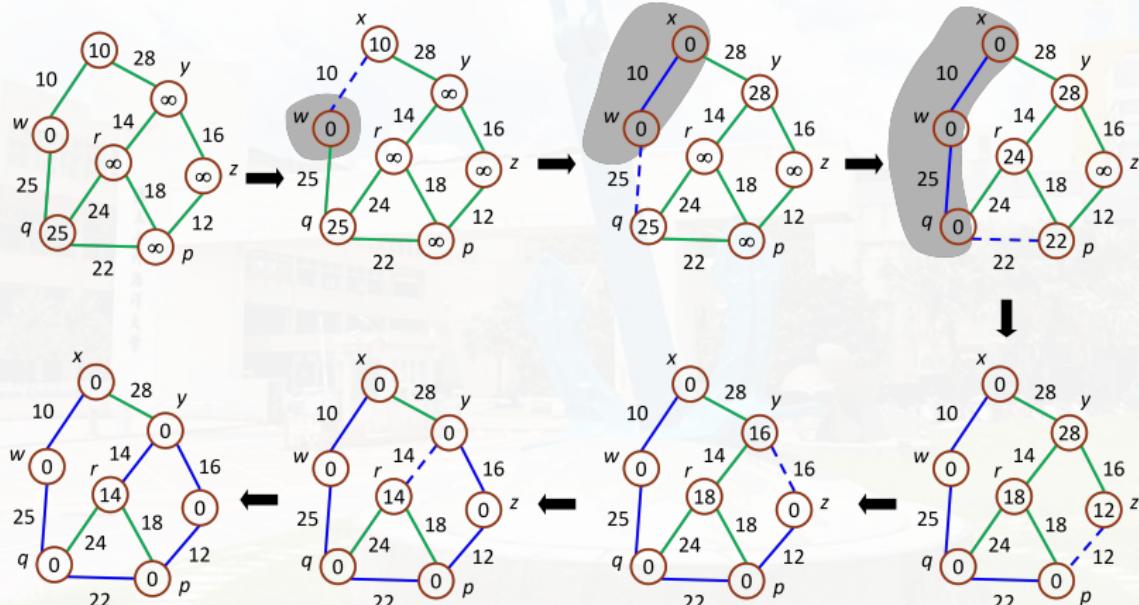
- Prim's algorithm begins with a tree T that contains **one arbitrary single vertex**.
- Next, we add a **least cost edge** (u, v) to T such that $T \cup (u, v)$ is **also a tree**.
- Repeat this edge addition until T contains $n - 1$ edges.

Prim's Algorithm (2/3)

- Prim's algorithm begins with a tree T that contains **one arbitrary single vertex**.
- Next, we add a **least cost edge** (u, v) to T such that $T \cup (u, v)$ is **also a tree**.
- Repeat this edge addition until T contains $n - 1$ edges.

Time complexity: $O(n^2)$.

Illustration of Prim's Algorithm

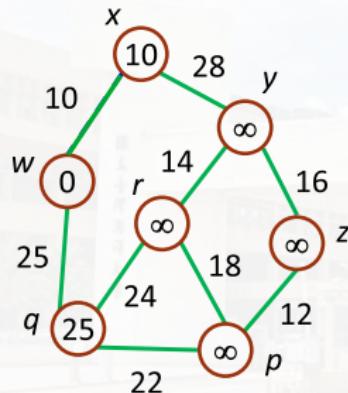


Prim's Algorithm (3/3)

```
T = {};
TV = {0};

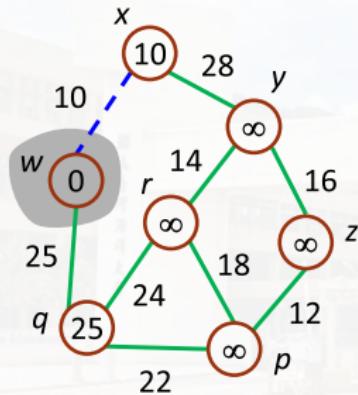
while (T contains fewer than n-1 edges) {
    let (u,v) be a least cost edge such that u is in TV and
    v is not in TV;
    if (there is no such edge )
        break;
    add v to TV;
    add (u,v) to T;
}
if (T contains fewer than n-1 edges)
    printf("No spanning tree\n");
```

- Each vertex $v \notin TV$ has a companion vertex “ $\text{near}(v)$ ” such that $\text{near}(v) \in TV$ and $\text{cost}(\text{near}(v), v)$ is minimum over all such choices for $\text{near}(v)$.
- Therefore, it takes $O(n)$ time to choose an edge.
- We can implement Prim's algorithm in $O(n^2)$ time.



	inMST	cost
w	F	∞
x	F	∞
y	F	∞
z	F	∞
p	F	∞
q	F	∞
r	F	∞

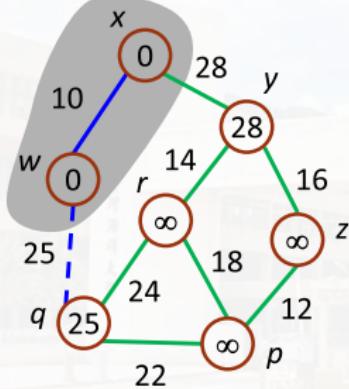
TV = { }



	inMST	cost
w	T	∞
x	F	10
y	F	∞
z	F	∞
p	F	∞
q	F	25
r	F	∞



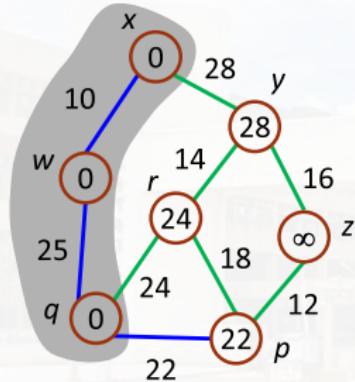
$TV = \{ w \}$



	inMST	cost
w	T	∞
x	T	10
y	F	28
z	F	∞
p	F	∞
q	F	25
r	F	∞



TV = { w, x }



	inMST	cost
w	T	∞
x	T	10
y	F	28
z	F	∞
p	F	22
q	T	25
r	F	24



$TV = \{ w, x, q \}$

Outline

- 1 Introduction
- 2 Kruskal's algorithm
- 3 Prim's algorithm
- 4 Sollin's/ Borůvka's algorithm
 - Supplementary (Cut Property)

Sollin's/ Borůvka's Algorithm (1/2)

- Sollin's algorithm selects several edges for inclusion in T at each stage.
- At the start of a stage, the selected edges, together with all the n vertices, form a spanning forest.
- During each stage, we select one edge for each tree in the forest.
 - This edge is a minimum cost edge that has exactly one vertex in the tree.

Sollin's/ Borůvka's Algorithm (2/2)

- **Note:** two trees in the forest could select the same edge.

Sollin's/ Borůvka's Algorithm (2/2)

- **Note:** two trees in the forest could select the same edge.
 - We need to **eliminate duplicate edges**.

Sollin's/ Borůvka's Algorithm (2/2)

- **Note:** two trees in the forest could select the same edge.
 - We need to **eliminate duplicate edges**.
- At the start of the first stage the set of selected edges is empty.

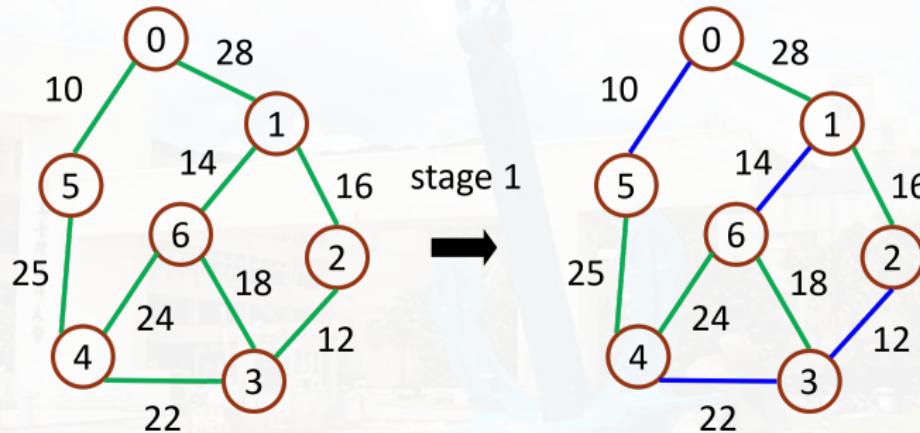
Sollin's/ Borůvka's Algorithm (2/2)

- **Note:** two trees in the forest could select the same edge.
 - We need to **eliminate duplicate edges**.
- At the start of the first stage the set of selected edges is empty.
- The algorithm terminates when there is **only one tree at the end of a stage** or **no edges remain for selection**.

We can implement Sollin's algorithm in $O(e \cdot \alpha(n) \log v)$ time (WHY?).

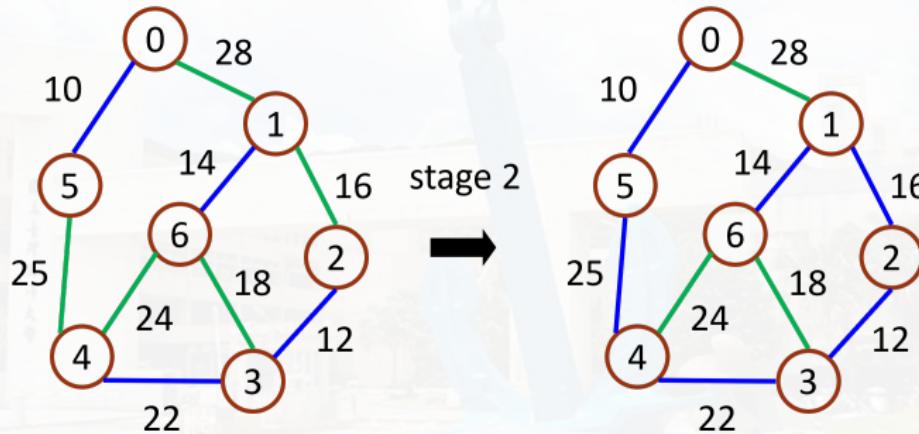
- $\alpha(n)$: the inverse Ackermann function.

Illustration of Sollin's/ Borůvka's Algorithm



- Stage 1: $(0, 5), (1, 6), (2, 3), (3, 2), (4, 3), (5, 0)$, and $(6, 1)$ are selected, and then duplicates are removed.

Illustration of Sollin's/ Borůvka's Algorithm



- Stage 2: $(5, 4)$, $(1, 2)$, and $(2, 1)$ are selected, and then duplicates are removed.

Borůvka-Sollin-MST (Use Union-Find structure to maintain components)

```

for each v in V: make_set(v) # initialization
MST = []
repeat until there is only one component: # stages
    for each component C:
        cheapest[C] = NULL # initialization
        for each edge (u,v) in E with weight w(u,v): # update the cheapest edges
            Cu = find_set(u); Cv = find_set(v); # find the trees u and v locate
            if Cu == Cv: continue # u and v are in the same tree, skipped
            if cheapest[Cu] = NULL or w(u,v) < w(cheapest[Cu]):
                cheapest[Cu] = (u,v)
            if cheapest[Cv] = NULL or w(u,v) < w(cheapest[Cv]):
                cheapest[Cv] = (u,v)
    for each component C:
        if cheapest[C] != NULL:
            (u,v) = cheapest[C]
            Cu = find_set(u); Cv = find_set(v);
            if Cu != Cv: # u and v belong two different components
                MST = MST + {(u,v)} # add (u,v) into MST
                union(Cu, Cv) # combine the two trees
return MST

```



Complexity: Number of Stages

Key observation

In each phase, every component selects at least one outgoing edge.

- Suppose at the beginning of a stage, we have k components.
- Each component chooses at least one edge to another component.
- Components merge in the best case by “pairing-up”, so the number of components drops to at most $\lceil \frac{k}{2} \rceil$.
- Starting from $k_0 = n$ components, after t stages we have at most $n/2^t$ components.
- We stop when the number of components reaches 1, which happens after $t = O(\lg n)$ stages.



Total running time

- There are $O(\log n)$ stages.
- Each stage has cost $O(e\alpha(n))$.
- Total time complexity of $O(e\alpha(n) \log n)$.
- Since $\alpha(n) \leq 4$ for all practical n , this is often informally written as $O(e \log n)$.

Notes on $\alpha(n)$

$\alpha(n) = O(\lg^* n)$, where $\lg^* n := \min\{k \geq 0 : \underbrace{\lg \lg \dots \lg}_{k \text{ times}}(n) \leq 1\}$.



Key to the Correctness

Cut Property (Key Tool for MST Algorithms)

A *cut* of G is a partition of the vertex set into two nonempty parts:

$$(S, V \setminus S), \quad \emptyset \subset S \subset V.$$

An edge $e = (u, v)$ *crosses* the cut if $u \in S$ and $v \in V \setminus S$ (or vice versa).

Cut Property

Let $(S, V \setminus S)$ be any cut, and let

$$e^* = \arg \min \{ w(e) : e \text{ crosses the cut} \}$$

be the (unique) minimum-weight edge crossing the cut. Then e^* belongs to every MST of G .

Outline

- 1 Introduction
- 2 Kruskal's algorithm
- 3 Prim's algorithm
- 4 Sollin's/ Borůvka's algorithm
 - Supplementary (Cut Property)

Cut Property (Stronger Version with Distinct Weights)

Distinct Edge Weights

Suppose all edge weights are distinct. Then each cut $(S, V \setminus S)$ has a *unique* minimum-weight edge crossing it.

Cut Property (Unique-Minimum Version)

Assume all edge weights are distinct. Let $(S, V \setminus S)$ be any cut, and let e be the unique minimum-weight edge crossing that cut. Then:

Every MST T^* of G contains e .



Exchange Argument: Setup

Setup

- Let $G = (V, E)$ be a connected, weighted graph.
- Fix a cut $(S, V \setminus S)$.
- Let $e = (u, v)$ be a minimum-weight edge crossing this cut:

$$u \in S, \quad v \in V \setminus S,$$

and $w(e) \leq w(f)$ for every edge f crossing the cut.

- Let T be an MST of G (an arbitrary one).

Goal

Show: there exists an MST T' such that $e \in T'$.

Exchange Argument: Adding e and Creating a Cycle

Case 1: $e \in T$

If e already belongs to T , then T itself is an MST containing e . We are done.

Case 2: $e \notin T$

- Consider the graph $T \cup \{e\}$.
- Since T is a tree, adding one edge e creates exactly one simple cycle C .
- This cycle C contains e , plus a unique simple path in T connecting u and v .

Key Observation

The cycle C must contain at least one other edge $f \in T$ that also crosses the cut $(S, V \setminus S)$.

Why the Cycle Contains Another Crossing Edge

- Edge $e = (u, v)$ crosses the cut:

$$u \in S, \quad v \in V \setminus S.$$

- Walk around the cycle C , starting from u .
 - Each time we traverse an edge that crosses the cut, we move from S to $V \setminus S$ or vice versa.
 - To return to u (which is in S), the number of crossings of the cut along the cycle must be even.
- Since e is one crossing edge in C , there must be at least one *other* edge f in C that crosses the same cut.
- This edge f belongs to T (because $C \subseteq T \cup \{e\}$ and $f \neq e$).

Conclusion

There exists $f \in T$ on the cycle C such that f crosses $(S, V \setminus S)$ and $f \neq e$.

Exchange Argument: Replacing f by e

Weight Comparison

- Both e and f cross the same cut $(S, V \setminus S)$.
- e is chosen as an edge of minimum-weight crossing this cut.
- Therefore, $w(e) \leq w(f)$.

Construct a New Spanning Tree

Define $T' = T - \{f\} \cup \{e\}$.

- Removing f breaks the cycle C , so T' is acyclic.
- Adding e keeps the graph connected (it reconnects the two components formed by removing f).
- Thus T' is a spanning tree.

Its total weight: $w(T') = w(T) - w(f) + w(e) \leq w(T)$.

Complexity: Cost per Stage and Overall

Cost per stage

- Use adjacency lists and a Union-Find structure.
- Initializing $\text{cheapest}[\cdot]$: $O(n)$.
- Scanning all edges $(u, v) \in E$:
 - Each scan does a constant number of find operations.
 - With weighted union (union-by-rank) and path compression, each find/union is $O(\alpha(n))$ amortized, where $\alpha(\cdot)$ is the inverse Ackermann function.
- Therefore, scanning all edges takes $O(e\alpha(n))$ time.
- Merging components with union also costs $O(n\alpha(n))$ in total per stage.
- Since $e \geq n - 1$ in a connected graph, the stage cost is $O(e\alpha(n))$.

Discussions

