

Cake Cutting Algorithms

Fair Allocation of Divisible Goods

Joseph Chuang-Chieh Lin

Department of Computer Science & Engineering,
National Taiwan Ocean University

Fall 2026



You may find that cake cutting is not just a piece of cake.

- *Ariel D. Procaccia: Cake cutting: not just child's play. Communicationis of the ACM, 2013.*
- *Jeff Edmonds & Kirk Pruhs: Cake cutting really is not a piece of cake. TALG 2011.*



Outline

- 1 Motivation and Model
- 2 Classic Cake Cutting Algorithms
- 3 Query Complexity
- 4 Optimal Cake Cutting
- 5 Price of Fairness




Outline

- 1 Motivation and Model
- 2 Classic Cake Cutting Algorithms
- 3 Query Complexity
- 4 Optimal Cake Cutting
- 5 Price of Fairness



Why Cake Cutting?

- The “cake” is a heterogeneous divisible resource: land, time, bandwidth, advertising slots, or computational resources.
- Agents may value different parts of the cake differently.
- **Goal:** design protocols that give formal fairness guarantees using limited information.



Computational Social Choice Viewpoint


Two concerns interact

- ① **Normative:** Which fairness properties should the allocation satisfy?
 - ② **Algorithmic:** How many operations or queries are needed to find such an allocation?
- Cake cutting is a clean testbed for continuous-input fair division.
 - Cake cutting also exposes the trade-off between fairness and efficiency: enforcing fairness may reduce the maximum achievable social welfare.



Basic Setting

- Agents: $N = \{1, 2, \dots, n\}$.
- Cake: the interval $[0, 1]$.
- Agent i has a valuation function V_i over intervals.
- We write $V_i(x, y)$ for $V_i([x, y])$.



$V_i(x, y)$ is agent i 's value for this interval.

Assumptions on Valuations

For every agent i , the valuation V_i satisfies:

- **Normalization:** $V_i(0, 1) = 1$.
- **Divisibility:** for any $[x, y]$ and $0 \leq \lambda \leq 1$, there exists $z \in [x, y]$ with

$$V_i(x, z) = \lambda V_i(x, y).$$

- **Nonnegativity:** $V_i(I) \geq 0$ for every interval I .
- **Additivity:** if I and I' are disjoint, then


$$V_i(I \cup I') = V_i(I) + V_i(I').$$



Value Density Functions

Often V_i is induced by a density v_i :

$$V_i(X) = \int_{x \in X} v_i(x) dx, \quad \int_0^1 v_i(x) dx = 1.$$



Pieces and Allocations

Definition

A **piece** of cake is a finite union of disjoint intervals.

Definition

An **allocation** is $A = (A_1, A_2, \dots, A_n)$, where A_i is agent i 's piece, the pieces are disjoint, and their union is the whole cake.




Agent 2's piece is non-contiguous here.

Contiguous vs. Non-Contiguous Allocations


Contiguous

Each A_i is a single interval.



Non-contiguous

A piece may be a union of intervals.



Contiguity is natural in some applications, but it often makes algorithmic problems much harder.

Fairness Properties

Proportionality

Every agent receives at least its fair share:

$$V_i(A_i) \geq \frac{1}{n} \quad \forall i \in N.$$



Fairness Properties

Proportionality

Every agent receives at least its fair share:

$$V_i(A_i) \geq \frac{1}{n} \quad \forall i \in N.$$

Envy-freeness

No agent prefers another agent's piece:

$$V_i(A_i) \geq V_i(A_j) \quad \forall i, j \in N.$$



Fairness Properties

Proportionality

Every agent receives at least its fair share:

$$V_i(A_i) \geq \frac{1}{n} \quad \forall i \in N.$$

Envy-freeness

No agent prefers another agent's piece:

$$V_i(A_i) \geq V_i(A_j) \quad \forall i, j \in N.$$

Equitability

All agents assign the same value to their own pieces:

$$V_{\textcolor{brown}{i}}(A_{\textcolor{brown}{i}}) = V_{\textcolor{teal}{j}}(A_{\textcolor{teal}{j}}) \quad \forall i, j \in N.$$



Envy-Freeness Implies Proportionality

Claim 1

Every envy-free allocation is proportional.



Envy-Freeness Implies Proportionality

Claim 1

Every envy-free allocation is proportional.

Proof idea.

Fix agent i . Since the pieces partition the cake,

$$\sum_{j \in N} V_i(A_j) = V_i([0, 1]) = 1.$$

Hence there exists some j such that $V_i(A_j) \geq 1/n$. If the allocation is envy-free,

$$V_i(A_i) \geq V_i(A_j) \geq \frac{1}{n}.$$




For **Two** Agents: Proportionality = Envy-Freeness

For $n = 2$, if $V_i(A_i) \geq 1/2$, then by additivity

$$V_i(A_{3-i}) = 1 - V_i(A_i) \leq 1/2 \leq V_i(A_i).$$

So agent i does not envy the other agent.



If each side is worth at least $1/2$ to its owner, no envy is possible.

For **Three or More** Agents: The Notions Split


- Envy-freeness is **stronger** than proportionality.
- A proportional allocation can still generate envy.

Example pattern

For $n = 3$, agent 1 may have

$$V_1(A_1) = \frac{1}{3}, \quad V_1(A_2) = \frac{1}{2}, \quad v_1(A_3) = \frac{1}{6}.$$

For agent $j \in \{2, 3\}$, $v_j(A_1) = v_j(A_2) = v_j(A_3) = \frac{1}{3}$. Proportionality holds for agent 1, but envy-freeness fails.



Note: Equitability Does Not Imply EF or PROP

Example

	g_1	g_2
a_1	0.4	0.6
a_2	0.6	0.4


Allocate $A_1 = \{g_1\}$, $A_2 = \{g_2\}$.

- Equitable: since $v_1(A_1) = 0.4 = v_2(A_2)$.
- Not envy-free:

$$v_1(A_2) = 0.6 > 0.4 = v_1(A_1), \quad v_2(A_1) = 0.6 > 0.4 = v_2(A_2).$$

- Not proportional:

$$v_i(A_i) = 0.4 < \frac{v_i(M)}{2} = 0.5 \quad \text{for } i = 1, 2.$$



Outline

- 1 Motivation and Model
- 2 Classic Cake Cutting Algorithms**
- 3 Query Complexity
- 4 Optimal Cake Cutting
- 5 Price of Fairness



A Nonconstructive Benchmark

Theorem 1 (Alon, 1987)

If $V_1, V_2, \dots, V_n$ are induced by continuous density functions, the cake **can be cut** in $n^2 - n$ places and partitioned into n pieces $A_1, A_2, \dots, A_n$ such that

$$V_i(A_j) = \frac{1}{n} \quad \forall i, j \in N.$$

A Nonconstructive Benchmark

Theorem 1 (Alon, 1987)

If $V_1, V_2, \dots, V_n$ are induced by continuous density functions, the cake **can be cut** in $n^2 - n$ places and partitioned into n pieces $A_1, A_2, \dots, A_n$ such that

$$V_i(A_j) = \frac{1}{n} \quad \forall i, j \in N.$$

- A very strong existence benchmark: every agent values every piece equally



A Nonconstructive Benchmark

Theorem 1 (Alon, 1987)

If $V_1, V_2, \dots, V_n$ are induced by continuous density functions, the cake **can be cut** in $n^2 - n$ places and partitioned into n pieces $A_1, A_2, \dots, A_n$ such that

$$V_i(A_j) = \frac{1}{n} \quad \forall i, j \in N.$$

- A very strong existence benchmark: every agent values every piece equally (equitability &



A Nonconstructive Benchmark

Theorem 1 (Alon, 1987)

If $V_1, V_2, \dots, V_n$ are induced by continuous density functions, the cake **can be cut** in $n^2 - n$ places and partitioned into n pieces $A_1, A_2, \dots, A_n$ such that

$$V_i(A_j) = \frac{1}{n} \quad \forall i, j \in N.$$

- A very strong existence benchmark: every agent values every piece equally (equitability & envyfreeness)



A Nonconstructive Benchmark

Theorem 1 (Alon, 1987)

If $V_1, V_2, \dots, V_n$ are induced by continuous density functions, the cake **can be cut** in $n^2 - n$ places and partitioned into n pieces $A_1, A_2, \dots, A_n$ such that

$$V_i(A_j) = \frac{1}{n} \quad \forall i, j \in N.$$

- A very strong existence benchmark: every agent values every piece equally (equitability & envyfreeness & proportionality).



A Nonconstructive Benchmark

Theorem 1 (Alon, 1987)

If $V_1, V_2, \dots, V_n$ are induced by continuous density functions, the cake **can be cut** in $n^2 - n$ places and partitioned into n pieces $A_1, A_2, \dots, A_n$ such that

$$V_i(A_j) = \frac{1}{n} \quad \forall i, j \in N.$$

- A very strong existence benchmark: every agent values every piece equally (equitability & envyfreeness & proportionality).
- We focus on **constructive** algorithms: how can we actually find fair allocations?



A Nonconstructive Benchmark

Theorem 1 (Alon, 1987)

If $V_1, V_2, \dots, V_n$ are induced by continuous density functions, the cake **can be cut** in $n^2 - n$ places and partitioned into n pieces $A_1, A_2, \dots, A_n$ such that


$$V_i(A_j) = \frac{1}{n} \quad \forall i, j \in N.$$

- A very strong existence benchmark: every agent values every piece equally (equitability & envyfreeness & proportionality).
- We focus on **constructive** algorithms: how can we actually find fair allocations?
 - Achievable but it requires $n^2 - n$ **continuous** operations [Brams et al. 2006] to obtain $V_i(A_j) = \frac{1}{n}$ we need to choose cut points with arbitrary real precision (e.g., moving-knife procedures, which lie outside the discrete que




Cut and Choose: $n = 2$

- 1 Agent 1 cuts the cake into two pieces it values equally.
- 2 Agent 2 chooses the preferred piece.
- 3 Agent 1 gets the remaining piece.



Cut and Choose: $n = 2$

- 1 Agent 1 cuts the cake into two pieces it values equally.
- 2 Agent 2 chooses the preferred piece.
- 3 Agent 1 gets the remaining piece.



Proportional, and for two agents therefore envy-free.

Exercise: Must the Cutter Split Equally?

Consider the standard cut-and-choose procedure for two agents.

- 1 Suppose agent 1 does not know agent 2's valuation. If agent 1 wants to maximize her guaranteed utility, how should she cut the cake?



Exercise: Must the Cutter Split Equally?

Consider the standard cut-and-choose procedure for two agents.

- 1 Suppose agent 1 does not know agent 2's valuation. If agent 1 wants to maximize her guaranteed utility, how should she cut the cake?
- 2 Formulate agent 1's decision as

$$\max_{(X_1, X_2)} \min\{V_1(X_1), V_1(X_2)\}.$$

What is the optimal guaranteed value?

Exercise: Must the Cutter Split Equally?

Consider the standard cut-and-choose procedure for two agents.

- 1 Suppose agent 1 does not know agent 2's valuation. If agent 1 wants to maximize her guaranteed utility, how should she cut the cake?
- 2 Formulate agent 1's decision as

$$\max_{(X_1, X_2)} \min\{V_1(X_1), V_1(X_2)\}.$$

What is the optimal guaranteed value?

- 3 Now suppose agent 1 knows agent 2's valuation. Must a utility-maximizing agent 1 still make a 50–50 cut according to V_1 ?



Exercise: Must the Cutter Split Equally?

Consider the standard cut-and-choose procedure for two agents.

- 1 Suppose agent 1 does not know agent 2's valuation. If agent 1 wants to maximize her guaranteed utility, how should she cut the cake?
- 2 Formulate agent 1's decision as

$$\max_{(X_1, X_2)} \min\{V_1(X_1), V_1(X_2)\}.$$

What is the optimal guaranteed value?

- 3 Now suppose agent 1 knows agent 2's valuation. Must a utility-maximizing agent 1 still make a 50–50 cut according to V_1 ?
- 4 Give an example in which agent 1 benefits from making an unequal cut because she knows which piece agent 2 will choose.



Dubins–Spanier (1961): Proportionality for Any n

- 1 Each remaining agent marks a point making the left prefix worth $1/n$ of the original cake.
- 2 The leftmost mark wins that prefix and exits.
- 3 Repeat with the remaining cake and agents.

The leftmost mark determines the prefix to be allocated.


remaining cake



Dubins–Spanier (1961): Proportionality for Any n

- 1 Each remaining agent marks a point making the left prefix worth $1/n$ of the original cake.
- 2 The leftmost mark wins that prefix and exits.
- 3 Repeat with the remaining cake and agents.


The leftmost mark determines the prefix to be allocated.



Dubins–Spanier (1961): Proportionality for Any n

- 1 Each remaining agent marks a point making the left prefix worth $1/n$ of the original cake.
- 2 The leftmost mark wins that prefix and exits.
- 3 Repeat with the remaining cake and agents.


The leftmost mark determines the prefix to be allocated.



Dubins–Spanier (1961): Proportionality for Any n

- 1 Each remaining agent marks a point making the left prefix worth $1/n$ of the original cake.
- 2 The leftmost mark wins that prefix and exits.
- 3 Repeat with the remaining cake and agents.


The leftmost mark determines the prefix to be allocated.



Dubins–Spanier (1961): Proportionality for Any n

- 1 Each remaining agent marks a point making the left prefix worth $1/n$ of the original cake.
- 2 The leftmost mark wins that prefix and exits.
- 3 Repeat with the remaining cake and agents.


The leftmost mark determines the prefix to be allocated.



Dubins–Spanier (1961): Proportionality for Any n

- 1 Each remaining agent marks a point making the left prefix worth $1/n$ of the original cake.
- 2 The leftmost mark wins that prefix and exits.
- 3 Repeat with the remaining cake and agents.

The leftmost mark determines the prefix to be allocated.



Why Dubins–Spanier Is Proportional?

Suppose agent i^* has the leftmost mark and receives $[0, x_{i^*}]$.

- By construction, $V_{i^*}(0, x_{i^*}) = 1/n$.
- For any other remaining agent j , leftmost means $x_{i^*} \leq x_j$, so

$$V_j(0, x_{i^*}) \leq V_j(0, x_j) = 1/n.$$

- Thus no remaining agent loses more than $1/n$ value when this piece is removed.

Inductively, everyone gets value at least $1/n$ of the original cake.



Why Dubins–Spanier Is Proportional?

Suppose agent i^* has the leftmost mark and receives $[0, x_{i^*}]$.

- By construction, $V_{i^*}(0, x_{i^*}) = 1/n$.
- For any other remaining agent j , leftmost means $x_{i^*} \leq x_j$, so

$$V_j(0, x_{i^*}) \leq V_j(0, x_j) = 1/n.$$


- Thus no remaining agent loses more than $1/n$ value when this piece is removed.

Inductively, everyone gets value at least $1/n$ of the original cake.

- For the last agent j , $V_j(A_j) \geq 1 - (n-1)/n = 1/n$.



Dubins–Spanier Is Not Envy-Free



agents exiting later may receive pieces
that earlier agents prefer

- Earlier exiting agents secure value $1/n$.
- An agent would never envy earlier agents, but they may still envy a later agent's piece.



Divide the cake into three regions:

$$I_1 = \left[0, \frac{1}{4}\right], \quad I_2 = \left[\frac{1}{4}, \frac{1}{2}\right], \quad I_3 = \left[\frac{1}{2}, 1\right].$$

	I_1	I_2	I_3
1	$\frac{1}{3}$	$\frac{1}{6}$	$\frac{1}{2}$
2	$\frac{1}{6}$	$\frac{1}{3}$	$\frac{1}{2}$
3	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{2}{3}$

- Round 1: agent 1 has the leftmost $\frac{1}{3}$ -mark, so $A_1 = I_1$.

Divide the cake into three regions:

$$I_1 = \left[0, \frac{1}{4}\right], \quad I_2 = \left[\frac{1}{4}, \frac{1}{2}\right], \quad I_3 = \left[\frac{1}{2}, 1\right].$$

	I_1	I_2	I_3
1	$\frac{1}{3}$	$\frac{1}{6}$	$\frac{1}{2}$
2	$\frac{1}{6}$	$\frac{1}{3}$	$\frac{1}{2}$
3	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{2}{3}$

- Round 1: agent 1 has the leftmost $\frac{1}{3}$ -mark, so $A_1 = I_1$.
- Round 2: agent 2 has the leftmost $\frac{1}{3}$ -mark, so $A_2 = I_2$.

Divide the cake into three regions:

$$I_1 = \left[0, \frac{1}{4}\right], \quad I_2 = \left[\frac{1}{4}, \frac{1}{2}\right], \quad I_3 = \left[\frac{1}{2}, 1\right].$$

	I_1	I_2	I_3
1	$\frac{1}{3}$	$\frac{1}{6}$	$\frac{1}{2}$
2	$\frac{1}{6}$	$\frac{1}{3}$	$\frac{1}{2}$
3	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{2}{3}$

- Round 1: agent 1 has the leftmost $\frac{1}{3}$ -mark, so $A_1 = I_1$.
- Round 2: agent 2 has the leftmost $\frac{1}{3}$ -mark, so $A_2 = I_2$.
- Agent 3 receives the remainder:

Divide the cake into three regions:

$$I_1 = \left[0, \frac{1}{4}\right], \quad I_2 = \left[\frac{1}{4}, \frac{1}{2}\right], \quad I_3 = \left[\frac{1}{2}, 1\right].$$

	I_1	I_2	I_3
1	$\frac{1}{3}$	$\frac{1}{6}$	$\frac{1}{2}$
2	$\frac{1}{6}$	$\frac{1}{3}$	$\frac{1}{2}$
3	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{2}{3}$

- Round 1: agent 1 has the leftmost $\frac{1}{3}$ -mark, so $A_1 = I_1$.
- Round 2: agent 2 has the leftmost $\frac{1}{3}$ -mark, so $A_2 = I_2$.
- Agent 3 receives the remainder:
 $A_3 = I_3$.




Divide the cake into three regions:

$$I_1 = \left[0, \frac{1}{4}\right], \quad I_2 = \left[\frac{1}{4}, \frac{1}{2}\right], \quad I_3 = \left[\frac{1}{2}, 1\right].$$

	I_1	I_2	I_3
1	$\frac{1}{3}$	$\frac{1}{6}$	$\frac{1}{2}$
2	$\frac{1}{6}$	$\frac{1}{3}$	$\frac{1}{2}$
3	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{2}{3}$

- Round 1: agent 1 has the leftmost $\frac{1}{3}$ -mark, so $A_1 = I_1$.
- Round 2: agent 2 has the leftmost $\frac{1}{3}$ -mark, so $A_2 = I_2$.
- Agent 3 receives the remainder: $A_3 = I_3$.




agent 1 envies agent 3

Divide the cake into three regions:

$$I_1 = \left[0, \frac{1}{4}\right], \quad I_2 = \left[\frac{1}{4}, \frac{1}{2}\right], \quad I_3 = \left[\frac{1}{2}, 1\right].$$

	I_1	I_2	I_3
1	$\frac{1}{3}$	$\frac{1}{6}$	$\frac{1}{2}$
2	$\frac{1}{6}$	$\frac{1}{3}$	$\frac{1}{2}$
3	$\frac{1}{6}$	$\frac{1}{6}$	$\frac{2}{3}$

- Round 1: agent 1 has the leftmost $\frac{1}{3}$ -mark, so $A_1 = I_1$.
- Round 2: agent 2 has the leftmost $\frac{1}{3}$ -mark, so $A_2 = I_2$.
- Agent 3 receives the remainder: $A_3 = I_3$.



agent 1 envies agent 3

$$V_1(A_1) = \frac{1}{3} < \frac{1}{2} = V_1(A_3).$$



Even–Paz (1984): Divide-and-Conquer

For simplicity assume n is a power of 2.


- 1 Given agents S and interval $[y, z]$, ask each $i \in S$ to mark x_i with

$$V_i(y, x_i) = \frac{1}{2} V_i(y, z).$$

- 2 Sort the marks from left to right.
- 3 Recurse on the left half of agents with the left subcake, and on the right half with the right subcake.




Even-Paz: One Recursive Split



The key invariant: each group of agents receives a subcake that each of them values sufficiently highly.

Even-Paz Recursion Tree



At depth k : each active agent values its current subcake at least $1/2^k$. At depth $\lg n$: each leaf agent gets value at least $1/n$.

Even–Paz: Why Proportional?

At recursive depth k :

- each active group has $n/2^k$ agents;
- every agent in that group values the group's cake at least $1/2^k$.

At depth $\lg n$, each group has one agent, and that agent receives a piece worth at least

$$\frac{1}{2^{\lg n}} = \frac{1}{n}.$$

Thus the allocation is proportional.




Selfridge–Conway: Envy-Freeness for $n = 3$

- First bounded envy-free algorithm for three agents.
 - Bounded number of cake-cutting operations.
- Core idea: trim the most valuable piece for agent 2 so it ties with the second-best piece.
- Allocate the main cake carefully, then allocate the trimming in a compensating order.




Selfridge–Conway: Initialization

- ① Agent 1 cuts the cake into three pieces it values equally: X_1, X_2, X_3 .



Selfridge–Conway: Initialization

- ① Agent 1 cuts the cake into three pieces it values equally: X_1, X_2, X_3 .
- ② Agent 2 trims its favorite piece (say X_1) so that it ties with its 2nd favorite.




Agent 2 trims her favorite piece X_1 so that $V_2(X_1 \setminus T) = V_2(X_2)$, assuming X_2 is her second-best piece.

Now the trimmed piece and the 2nd-best piece are tied according to agent 2.



Selfridge–Conway: Allocate the Main Cake

- ③ Agent 3 chooses first from the three main pieces.
- ④ If agent 3 did not choose the trimmed piece, agent 2 gets the trimmed piece (i.e., $X_1 \setminus T$ here).
- ⑤ Otherwise agent 2 chooses from the two remaining pieces.
- ⑥ Agent 1 receives the last main piece.



Selfridge–Conway: Allocate the Trimming

Let a be the agent among $\{2, 3\}$ who received the trimmed main piece, and let a' be the other.

- ⑥ Agent a' cuts the trimming into three equal-value pieces.
- ⑦ Selection order: a , then agent 1, then a' .



Selfridge–Conway: Proof Sketch

- Main cake is envy-free:
 - agent 3 chooses first;
 - agent 2 receives a piece tied for best in its valuation;
 - agent 1 receives an untrimmed piece, which it values as one of the original thirds.
- Trimming allocation protects a and a' by the cut-and-choose logic.
- Agent 1 cannot use the trimming to envy a overall: even all the trimming would only reconstruct one original piece for agent 1.



Classic Algorithms: Comparison

Algorithm	Agents	Guarantee	Pieces
Cut and choose	2	Prop. + EF	contiguous
Dubins–Spanier	any n	Prop.	contiguous
Even–Paz	any n	Prop.	contiguous
Selfridge–Conway	3	EF	≤ 2 intervals per agent

EF = envy-free; Prop. = proportional.



Outline

- 1 Motivation and Model
- 2 Classic Cake Cutting Algorithms
- 3 Query Complexity**
- 4 Optimal Cake Cutting
- 5 Price of Fairness



Why a New Complexity Model?

Standard input-size analysis is not directly applicable:

- a valuation function over $[0, 1]$ may not have a finite discrete representation;
- algorithms often interact with agents through cuts and value reports;
- complexity is measured by the number of permitted queries.

Question

How many valuation queries are necessary and sufficient to guarantee **fairness**?



The Robertson–Webb Model

Two query types are allowed.

Evaluation query

$$\text{eval}_i(x, y) = V_i(x, y).$$

Agent i reports its **value** for interval $[x, y]$.



Cut query

$$\text{cut}_i(x, \alpha) = y,$$

where y is the **leftmost point** such that $V_i(x, y) = \alpha$.




Query Types



Simulating Cut and Choose

- Ask agent 1: $\text{cut}_1(0, 1/2) = y$.
- Ask agent 2: $\text{eval}_2(0, y)$.
- If agent 2 values $[0, y]$ at least $1/2$, give it $[0, y]$; otherwise give it $[y, 1]$.



Query Complexity: Dubins–Spanier

- At a stage with k remaining agents, ask each remaining agent one cut query.

$$k = n, n - 1, \dots, 2.$$




Query Complexity: Dubins–Spanier

- At a stage with k remaining agents, ask each remaining agent one cut query.

$$k = n, n - 1, \dots, 2.$$

- Thus the number of queries is

$$\sum_{k=2}^n k = \Theta(n^2).$$



Query Complexity: Even–Paz

At each recursive level, every active agent makes one cut query.

- Level 0: n queries.
- Level 1: $2 \cdot (n/2) = n$ queries.
- Level k : $2^k \cdot (n/2^k) = n$ queries.

There are $\lg n$ levels, so the total is

$$n \lg n = \Theta(n \log n).$$



Optimality of Even–Paz

Theorem 2 (Edmonds–Pruhs, SODA 2006)

Any proportional cake-cutting algorithm requires

$$\Omega(n \log n)$$

queries in the Robertson–Webb model.

Consequence

Even–Paz is **asymptotically optimal** for proportional cake cutting in this query model.



Lower Bound Strategy: Thin-Rich Pieces

We separate the problem of finding a proportional allocation into problems that must be solved **for each agent individually**.

Fix an agent i .

A piece X is:

- **thin** if $\ell(X) \leq 2/n$; ($\ell(X)$: length of the cake piece X)
- **rich** if $V_i(X) \geq 1/n$;
- **thin-rich** if both conditions hold.

Lemma 1

If **finding a thin-rich piece needs $T(n)$ queries**, then proportional cake cutting needs $\Omega(nT(n))$ queries.

Sketch of the proof of Lemma 1

- Each query only acquires information about the valuation function of a single agent, hence the interaction with one agent cannot help us find a thin-rich piece with respect to another agent.
- In any **proportional** allocation, every allocated piece **must be rich for its owner**.
- At most $n/2$ allocated pieces can have length more than $2/n$.
- Therefore, **in any proportional allocation, at least $n/2$ agents receive thin-rich pieces.**




Sketch of the proof of Lemma 1

- Each query only acquires information about the valuation function of a single agent, hence the interaction with one agent cannot help us find a thin-rich piece with respect to another agent.
- In any **proportional** allocation, every allocated piece **must be rich for its owner**.
- At most $n/2$ allocated pieces can have length more than $2/n$.
- Therefore, **in any proportional allocation, at least $n/2$ agents receive thin-rich pieces.** $\Rightarrow$ one needs to solve the thin-rich problem w.r.t. $\geq n/2$ agents.



Valuation Function: a Value Tree

- The lower bound uses **adversarial** valuation functions represented by a ternary tree.
- An internal node: the union of the intervals at the leaves of the subtree rooted at that node,



From Thin-Rich Pieces to Thin-Rich Leaves

Lemma 2

Suppose the valuation function is derived from a value tree.

If a **thin-rich piece** can be found using $T(n)$ Robertson–Webb queries, then a **thin-rich leaf** of the value tree can be found using $O(T(n))$ queries.



From Thin-Rich Pieces to Thin-Rich Leaves

Lemma 2

Suppose the valuation function is derived from a value tree.

If a **thin-rich piece** can be found using $T(n)$ Robertson–Webb queries, then a **thin-rich leaf** of the value tree can be found using $O(T(n))$ queries.

If a **thin-rich leaf** of the value tree cannot be found in $O(T(n))$ queries, then a **thin-rich piece** cannot be found in $T(n)$ queries.



From Thin-Rich Pieces to Thin-Rich Leaves

Lemma 2

Suppose the valuation function is derived from a value tree.

If a **thin-rich piece** can be found using $T(n)$ Robertson–Webb queries, then a **thin-rich leaf** of the value tree can be found using $O(T(n))$ queries.

- A leaf is an interval of length $2/n$.
- A piece X is thin-rich if

$$\ell(X) \leq \frac{2}{n} \quad \text{and} \quad V(X) \geq \frac{1}{n}.$$

- **Claim:** It is enough to prove lower bounds for finding a rich leaf.



Key Observation: Some Interval Has High Density

- Suppose an algorithm finds a thin-rich piece X .
- Write X as a union of intervals. Since X is rich, $V(X) \geq \frac{1}{n}$.



Key Observation: Some Interval Has High Density

- Suppose an algorithm finds a thin-rich piece X .
- Write X as a union of intervals. Since X is rich, $V(X) \geq \frac{1}{n}$.

Claim: The density of I^* is $\geq 1/2$

There exists an interval $I^* \subseteq X$ such that $V(I^*) \geq \frac{\ell(I^*)}{2}$.



Key Observation: Some Interval Has High Density

- Suppose an algorithm finds a thin-rich piece X .
- Write X as a union of intervals. Since X is rich, $V(X) \geq \frac{1}{n}$.

Claim: The density of I^* is $\geq 1/2$

There exists an interval $I^* \subseteq X$ such that $V(I^*) \geq \frac{\ell(I^*)}{2}$.

- Otherwise, if every interval $I \subseteq X$ satisfied

$$V(I) < \frac{\ell(I)}{2},$$

then

$$V(X) = \sum_{I \subseteq X} V(I) < \sum_{I \subseteq X} \frac{\ell(I)}{2} = \frac{\ell(X)}{2} \leq \frac{1}{n},$$

contradicting that X is rich.



Finding a Thin-Rich Leaf

- Since X is thin, $\ell(X) \leq \frac{2}{n}$.
- Hence the interval $I^* \subseteq X$ also has length at most $2/n$. Because each leaf has length $2/n$, we know that I^* intersects at most two leaves.
 - Note: I^* is not assumed to coincide with a leaf or to be aligned with leaf boundaries.



Finding a Thin-Rich Leaf

- Since X is thin, $\ell(X) \leq \frac{2}{n}$.
- Hence the interval $I^* \subseteq X$ also has length at most $2/n$. Because each leaf has length $2/n$, we know that I^* intersects at most two leaves.
 - Note: I^* is not assumed to coincide with a leaf or to be aligned with leaf boundaries.
- The value density is constant on each leaf. Since the average density on I^* is at least $1/2$, one of the leaves it intersects, say u , has density at least $1/2$.
- Thus, $V(u) \geq \frac{\ell(u)}{2} = \frac{1}{n}$. So, u is rich, and since $\ell(u) = 2/n$, it is also thin.



Finding a Thin-Rich Leaf

- Since X is thin, $\ell(X) \leq \frac{2}{n}$.
- Hence the interval $I^* \subseteq X$ also has length at most $2/n$. Because each leaf has length $2/n$, we know that I^* intersects at most two leaves.
 - Note: I^* is not assumed to coincide with a leaf or to be aligned with leaf boundaries.
- The value density is constant on each leaf. Since the average density on I^* is at least $1/2$, one of the leaves it intersects, say u , has density at least $1/2$.
- Thus, $V(u) \geq \frac{\ell(u)}{2} = \frac{1}{n}$. So, u is rich, and since $\ell(u) = 2/n$, it is also thin.

Query cost

After $T(n)$ queries, the output piece X has only $O(T(n))$ relevant intervals. Using $O(T(n))$ additional eval queries, we can evaluate the leaves intersecting X and find such a leaf u .

Value Tree: Rich Leaves Require Heavy Edges

- Each leaf represents a cake interval of length $2/n$. Since the cake has total length 1, the value tree has $\frac{1}{2/n} = \frac{n}{2}$ leaves.
- Let the height be $H = \log_3(n/2)$. If a leaf u has $q(u)$ heavy edges on its root-to-leaf path, then

$$V(u) = \left(\frac{1}{2}\right)^{q(u)} \left(\frac{1}{4}\right)^{H-q(u)} = \left(\frac{1}{4}\right)^{H-\frac{q(u)}{2}}.$$



Value Tree: Rich Leaves Require Heavy Edges

- Each leaf represents a cake interval of length $2/n$. Since the cake has total length 1, the value tree has $\frac{1}{2/n} = \frac{n}{2}$ leaves.
- Let the height be $H = \log_3(n/2)$. If a leaf u has $q(u)$ heavy edges on its root-to-leaf path, then

$$\frac{1}{n} \leq V(u) = \left(\frac{1}{2}\right)^{q(u)} \left(\frac{1}{4}\right)^{H-q(u)} = \left(\frac{1}{4}\right)^{H-\frac{q(u)}{2}}.$$

- To be rich, $V(u) \geq 1/n$.



Value Tree: Rich Leaves Require Heavy Edges

- Each leaf represents a cake interval of length $2/n$. Since the cake has total length 1, the value tree has $\frac{1}{2/n} = \frac{n}{2}$ leaves.
- Let the height be $H = \log_3(n/2)$. If a leaf u has $q(u)$ heavy edges on its root-to-leaf path, then

$$\frac{1}{n} \leq V(u) = \left(\frac{1}{2}\right)^{q(u)} \left(\frac{1}{4}\right)^{H-q(u)} = \left(\frac{1}{4}\right)^{H-\frac{q(u)}{2}}.$$

- To be rich, $V(u) \geq 1/n$.

A rich leaf must have a constant fraction of heavy edges on its path.
Discovering that information costs $\Omega(\log n)$ queries.



By $2H - \lg n = \Omega(\lg n)$:

$$\left(\frac{1}{4}\right)^{H - \frac{q(u)}{2}} \geq \frac{1}{n}$$

$$\Rightarrow 4^{H - \frac{q(u)}{2}} \leq n$$

$$\Rightarrow 2^{2H - q(u)} \leq n$$

$$\Rightarrow q(u) = \Omega(\lg n).$$

Note about the queries

- What we are interested in is



Note about the queries

- What we are interested in is how to simulate **cut** and **eval** queries by revealing edges of the value tree.



Note about the queries

- What we are interested in is how to simulate **cut** and **eval** queries by **revealing edges of the value tree**.
- A node u is **revealed** if the weights of *all* edges of every node on the path from root to u are known to the algorithm.



From Tree Queries to Robertson–Webb Queries

The proof answers each Robertson–Webb query by revealing enough of the value tree.

- An $\text{eval}(x, y)$ query reveals **paths to the leaves** containing x and y .
- A $\text{cut}(x, \alpha)$ query reveals **paths** sufficient to identify **the cut point**.
- ★ **Maintaining the invariant: After k queries, every root-to-leaf path contains at most $2k$ edges known to be heavy.**
 - Because we have shown that $\Omega(\lg n)$ edges on the path must be known to be heavy, the theorem will directly follow.



From Tree Queries to Robertson–Webb Queries


The proof answers each Robertson–Webb query by revealing enough of the value tree.

- An $\text{eval}(x, y)$ query reveals **paths to the leaves** containing x and y .
- A $\text{cut}(x, \alpha)$ query reveals **paths** sufficient to identify **the cut point**.
- ★ **Maintaining the invariant: After k queries, every root-to-leaf path contains at most $2k$ edges known to be heavy.**
 - Because we have shown that $\Omega(\lg n)$ edges on the path must be known to be heavy, the theorem will directly follow.

$$c \lg n \leq q(u) \leq 2k \implies k \geq \frac{c}{2} \lg n = \Omega(\lg n).$$



Value-tree Illustration



Revealed Nodes: Useful Facts

- ★ Say V , a value function V_i of some agent i .
- If a node u is revealed, then its value $V(u)$ is known.
- If a revealed node corresponds to interval $u = [x, y]$, then $V(0, x)$, $V(x, y)$, $V(y, 1)$ are all known.
- If $z \in [x, y]$ and $[x, y]$ is a revealed leaf, then $V(0, z)$ can be computed because value density is uniform inside each leaf.
- If w lies in a revealed leaf and α is given, we can locate the leaf containing the point z satisfying $V(w, z) = \alpha$.



Why does the invariant hold? (1/2)

- Initially, the invariant holds trivially (no queries, $k = 0$, at most $2k = 0$ edges are known to be heavy)
- Assume that query $k + 1$ is an $\text{eval}(x, y)$ query.
 - Reveal the path ($r = u_0 \rightarrow u_1 \rightarrow \dots \rightarrow u_H = u$) to the leaf u containing x .
 - For every newly revealed node u_i , the edge $u_i \rightarrow u_{i+1}$ is declared **light**.
 - Other two child edges: one **light**, and one **heavy**.
- Repeat the same procedure for y .

increase of known heavy edges per query $\leq 1 + 1 = 2$.

Every root-to-leaf path now has $\leq 2k + 2 = 2(k + 1)$ known heavy edges.



Why does the invariant hold? (1/2)

- Initially, the invariant holds trivially (no queries, $k = 0$, at most $2k = 0$ edges are known to be heavy)
- Assume that query $k + 1$ is an $\text{eval}(x, y)$ query.
 - Reveal the path ($r = u_0 \rightarrow u_1 \rightarrow \dots \rightarrow u_H = u$) to the leaf u containing x .
 - For every newly revealed node u_i , the edge $u_i \rightarrow u_{i+1}$ is declared **light**.
 - Other two child edges: one **light**, and one **heavy**.
- Repeat the same procedure for y .

increase of known heavy edges per query $\leq 1 + 1 = 2$.

Every root-to-leaf path now has $\leq 2k + 2 = 2(k + 1)$ known heavy edges.

In a rooted tree, the intersection of any root-to-leaf path with the path to x is a prefix. Once the two paths diverge, they never meet again.



Why does the invariant hold? (2/2)

- Assume that query $k + 1$ is a $\text{cut}(w, \alpha)$. Let z be the point satisfying $V(w, z) = \alpha$.
 - Reveal the leaf containing w .
 - Find the least common ancestor u of the leaf containing w and the leaf containing z .
 - Reveal the path from u toward the leaf containing z .
- The adversary chooses (unrevealed) edge labels so that the path between u and z follows **light** edges.
- At a newly revealed node u' , let β be the remaining value needed from u' .
 - if $\beta/V(u') > 1/2 \Rightarrow$ make the left edge **heavy**;then the path must follow a middle/right **light** edge.
 - if $\beta/V(u') \leq 1/2 \Rightarrow$ make the right edge **heavy**;again the path follows a **light** edge.



Envy-Free Cake Cutting Is Harder

- Selfridge–Conway gives a bounded envy-free algorithm for $n = 3$.
- Brams–Taylor (1995) gives a *finite* envy-free algorithm for any n .
 - The algorithm is guaranteed to *terminate*.
- But for $n \geq 4$, the number of queries of Brams–Taylor is unbounded as a function of n alone.

Historical open problem

Is there a bounded envy-free algorithm for arbitrary n ?



Envy-Free Cake Cutting Is Harder

- Selfridge–Conway gives a bounded envy-free algorithm for $n = 3$.
- Brams–Taylor (1995) gives a *finite* envy-free algorithm for any n .
 - The algorithm is guaranteed to *terminate*.
- But for $n \geq 4$, the number of queries of Brams–Taylor is unbounded as a function of n alone.

Historical open problem

Is there a bounded envy-free algorithm for arbitrary n ?

Aziz–Mackenzie (STOC 2016) answered this positively: a bounded envy-free protocol exists for every n .



Bounded vs. Finite

Finite: $\forall n, \forall V, \quad Q(n, V) < \infty.$

Bounded: $\exists f : \mathbb{N} \rightarrow \mathbb{N}$ such that $\forall n, \forall V, \quad Q(n, V) \leq f(n).$



Known Lower Bounds for Envy-Freeness

Theorem (Stromquist, Electron. J. Comb. 2008)

For $n \geq 3$, no finite envy-free cake-cutting algorithm outputs contiguous allocations.

Theorem (Procaccia, AAI 2009; Theorem 3.1)

Any envy-free cake-cutting algorithm requires $\Omega(n^2)$ queries in the Robertson–Webb model.

This separates envy-freeness from proportionality:

$\Theta(n \log n)$ for proportionality vs. $\Omega(n^2)$ for envy-freeness.



Piecewise Uniform Valuations Do Not Avoid the Difficulty

Theorem (Kurokawa–Procaccia–Wang, AAAI 2013)

Let $\mathcal{A}$ be an envy-free algorithm for *piecewise uniform* valuations using at most $f(n)$ Robertson–Webb queries. Then $\mathcal{A}$ can compute an envy-free algorithm for **general** valuations using at most $f(n)$ queries.

Piecewise uniform valuations are expressive enough to capture the hardness of envy-free cake cutting.



ε -Envy-Freeness

Relax envy-freeness to

$$V_i(A_i) \geq V_i(A_j) - \varepsilon \quad \forall i, j.$$

- ① Ask each agent to cut the cake into $\lceil 1/\varepsilon \rceil$ disjoint intervals worth of value about ε .
- ② Sort all $O(n/\varepsilon)$ cut points.
- ③ Treat intervals between adjacent cut points as indivisible goods.
- ④ Allocate these small goods by round-robin selection.

This gives a simple and efficient approximate envy-free allocation.

- $O(n/\varepsilon)$ cut queries and $O(n^2/\varepsilon)$ eval queries.



Round-Robin Argument for ε -EF



- After an agent first picks, it picks before each other agent in every subsequent round.
- The only possible envy comes from goods selected earlier, each worth at most ε .

Round-Robin Argument for ε -EF



- After an agent first picks, it picks before each other agent in every subsequent round.
- The only possible envy comes from goods selected earlier, each worth at most ε .
- We will talk about this algorithm more for **fair division of indivisible goods**.



Outline


- 1 Motivation and Model
- 2 Classic Cake Cutting Algorithms
- 3 Query Complexity
- 4 Optimal Cake Cutting**
- 5 Price of Fairness



From Query Model to Explicit Representation


For valuations that are assumed to be explicitly represented:

- Example: piecewise constant densities.
- The input lists intervals and their density values.
- Now the computational task is to **optimize over allocations**, not to elicit valuations.
 - The notion of efficiency: **Social Welfare**.



Common Refinement of Segments

Collect all boundaries reported by all agents, plus 0 and 1. Let $\mathcal{J}$ be the set of intervals between consecutive marks.



Crucial property

For each $J \in \mathcal{J}$ and each agent i , the density v_i is constant on J .

Computing Equitable / Envy-Free Allocations

If each refined interval J is split into n equal-length subintervals and each agent gets one subinterval from every J , then every agent values every **final bundle** equally:

$$V_i(A_j) = \frac{1}{n} \quad \forall i, j.$$



This is constructive under explicit piecewise constant valuations.

Social Welfare

The **utilitarian social welfare** of allocation A is

$$\text{sw}(A) = \sum_{i \in N} V_i(A_i).$$

- Once welfare is introduced, fairness becomes a constraint in an optimization problem.

Optimization target

Maximize social welfare subject to proportionality, envy-freeness, or equitability.



LP for Optimal Proportional Allocation

Let f_i^J be the fraction of refined interval J allocated to agent i .

$$\begin{aligned} \max \quad & \sum_{i \in N} \sum_{J \in \mathcal{J}} f_i^J V_i(J) \\ \text{s.t.} \quad & \sum_{i \in N} f_i^J \leq 1 \quad \forall J \in \mathcal{J}, \\ & \sum_{J \in \mathcal{J}} f_i^J V_i(J) \geq \frac{1}{n} \quad \forall i \in N, \\ & f_i^J \geq 0 \quad \forall i, J. \end{aligned}$$

- A polynomial-size linear program for explicitly represented piecewise constant valuations.



LP Constraints for Other Fairness Notions

Using the same variables f_i^J , other constraints are linear.

Envy-freeness

$$\sum_J f_i^J V_i(J) \geq \sum_J f_j^J V_i(J) \quad \forall i, j.$$

Equitability

$$\sum_J f_i^J V_i(J) = \sum_J f_j^J V_j(J) \quad \forall i, j.$$


Thus LPs can compute optimal fair allocations under non-contiguous pieces.



Contiguity Makes Optimization Hard

Theorem (Bei et al., AAAI 2012)

*Given explicit piecewise constant valuations, if allocations must be both proportional and **contiguous**, then maximizing social welfare is NP-hard to obtain an $o(\sqrt{n})$ -approximation factor.*



Outline

- 1 Motivation and Model
- 2 Classic Cake Cutting Algorithms
- 3 Query Complexity
- 4 Optimal Cake Cutting
- 5 Price of Fairness



Price of Fairness: Definition

For a fairness property $\mathcal{F}$, define

$$\text{PoF}(\mathcal{F}) = \sup_{v_1, v_2, \dots, v_n} \frac{\max_A \text{sw}(A)}{\max_{A \in \mathcal{F}} \text{sw}(A)}.$$

- Numerator: unconstrained welfare optimum.
- Denominator: best welfare among fair allocations.
- Larger price means greater worst-case welfare loss due to fairness.



Price of Proportionality

Theorem (Caragiannis et al., WINE 2009)

The price of proportionality is

$$\Theta(\sqrt{n}).$$

Interpretation

A proportionality constraint can reduce optimal social welfare by a factor on the order of $\sqrt{n}$, and this bound is tight.



Upper Bound Idea for Proportionality

Let $\mathbf{A}^* = (A_1^*, A_2^*, \dots, A_n^*)$ be the unconstrained optimal allocation and define

$$L = \left\{ i : V_i(A_i^*) \geq \frac{1}{\sqrt{n}} \right\}, \quad S = N \setminus L.$$

- L : the set of “large” agents; S : the set of “small” agents;



Upper Bound Idea for Proportionality

Let $\mathbf{A}^* = (A_1^*, A_2^*, \dots, A_n^*)$ be the unconstrained optimal allocation and define

$$L = \left\{ i : V_i(A_i^*) \geq \frac{1}{\sqrt{n}} \right\}, \quad S = N \setminus L.$$

- L : the set of “large” agents; S : the set of “small” agents;
- Case 1: Many agents are “large”: $|L| \geq \sqrt{n}$. Then, we carefully redistribute pieces to protect everyone while losing at most a $\sqrt{n}$ factor.



Upper Bound Idea for Proportionality

Let $\mathbf{A}^* = (A_1^*, A_2^*, \dots, A_n^*)$ be the unconstrained optimal allocation and define

$$L = \left\{ i : V_i(A_i^*) \geq \frac{1}{\sqrt{n}} \right\}, \quad S = N \setminus L.$$

- L : the set of “large” agents; S : the set of “small” agents;
- Case 1: Many agents are “large”: $|L| \geq \sqrt{n}$. Then, we carefully redistribute pieces to protect everyone while losing at most a $\sqrt{n}$ factor.
- Case 2: Few agents are large: $|L| < \sqrt{n}$. Then, the unconstrained social welfare is already $O(\sqrt{n})$, while any proportional allocation has welfare at least 1.
 - $\text{sw}(\mathbf{A}^*) \leq |L| + |S|/\sqrt{n} < 2\sqrt{n}$, while for any proportional $\mathbf{A}$,
 $\text{sw}(\mathbf{A}) \geq \sum_{i \in N} 1/n = 1$.



Upper Bound Idea for Proportionality

Let $\mathbf{A}^* = (A_1^*, A_2^*, \dots, A_n^*)$ be the unconstrained optimal allocation and define

$$L = \left\{ i : V_i(A_i^*) \geq \frac{1}{\sqrt{n}} \right\}, \quad S = N \setminus L.$$

- L : the set of “large” agents; S : the set of “small” agents;
- Case 1: Many agents are “large”: $|L| \geq \sqrt{n}$. Then, we carefully redistribute pieces to protect everyone while losing at most a $\sqrt{n}$ factor.
- Case 2: Few agents are large: $|L| < \sqrt{n}$. Then, the unconstrained social welfare is already $O(\sqrt{n})$, while any proportional allocation has welfare at least 1.
 - $\text{sw}(\mathbf{A}^*) \leq |L| + |S|/\sqrt{n} < 2\sqrt{n}$, while for any proportional $\mathbf{A}$,
 $\text{sw}(\mathbf{A}) \geq \sum_{i \in N} 1/n = 1$.



Price of Proportionality: Case 1 Setup

Case 1

Assume

$$|L| \geq \sqrt{n}.$$

Then

$$|S| = n - |L| \leq n - \sqrt{n}, \quad \sqrt{n} + |S| \leq n.$$

- Agents in L get “large” value in the optimal allocation.
- We transform $\mathbf{A}^*$ into a **proportional allocation** $\mathbf{A}$.
- Goal:

$$V_i(A_i) \geq \frac{V_i(A_i^*)}{\sqrt{n}} \quad \forall i \in N.$$



Case 1: Redistributing the Pieces

We **redistribute** each optimal piece A_i^* separately.

Pieces originally owned by *small* agents

For each $i \in S$, divide A_i^* proportionally among all agents j in S :

$$V_j(A_j \cap A_i^*) \geq \frac{V_j(A_i^*)}{|S|} \quad \forall j \in S.$$

Case 1: Redistributing the Pieces

We **redistribute** each optimal piece A_i^* separately.

Pieces originally owned by *small* agents

For each $i \in S$, divide A_i^* proportionally among all agents j in S :

$$V_j(A_j \cap A_i^*) \geq \frac{V_j(A_i^*)}{|S|} \quad \forall j \in S.$$

Pieces originally owned by *large* agents

For each $i \in L$, divide A_i^* among $\{i\} \cup S$ by treating agent i as $\sqrt{n}$ identical copies. Then

$$V_i(A_i \cap A_i^*) \geq \frac{\sqrt{n}}{\sqrt{n} + |S|} V_i(A_i^*),$$

and for every $j \in S$,

$$V_j(A_j \cap A_i^*) \geq \frac{1}{\sqrt{n} + |S|} V_j(A_i^*).$$

Case 1: Why the loss is at most $\sqrt{n}$

First, $\mathbf{A}$ is proportional.

- For $i \in L$:

$$V_i(A_i) \geq \frac{\sqrt{n}}{\sqrt{n} + |S|} V_i(A_i^*) \geq \frac{1}{\sqrt{n} + |S|} \geq \frac{1}{n}.$$

- For $i \in S$:

$$V_i(A_i) \geq \sum_{k \in L} \frac{V_i(A_k^*)}{\sqrt{n} + |S|} + \sum_{j \in S} \frac{V_i(A_j^*)}{|S|} \geq \sum_{j \in N} \frac{V_i(A_j^*)}{n} = \frac{1}{n}.$$



Case 1: Why the loss is at most $\sqrt{n}$

Moreover,

- for $i \in L$:

$$V_i(A_i) \geq \frac{\sqrt{n}}{\sqrt{n} + |S|} V_i(A_i^*) \geq \frac{\sqrt{n}}{n} V_i(A_i^*) = \frac{1}{\sqrt{n}} V_i(A_i^*).$$

- For $i \in S$, $V_i(A_i^*) < \frac{1}{\sqrt{n}}$ (by definition), and $V_i(A_i) \geq \frac{1}{n}$ (by proportionality we just built).



Case 1: Why the loss is at most $\sqrt{n}$

Moreover,

- for $i \in L$:

$$V_i(A_i) \geq \frac{\sqrt{n}}{\sqrt{n} + |S|} V_i(A_i^*) \geq \frac{\sqrt{n}}{n} V_i(A_i^*) = \frac{1}{\sqrt{n}} V_i(A_i^*).$$

- For $i \in S$, $V_i(A_i^*) < \frac{1}{\sqrt{n}}$ (by definition), and $V_i(A_i) \geq \frac{1}{n}$ (by proportionality we just built). We have

Case 1: Why the loss is at most $\sqrt{n}$

Moreover,

- for $i \in L$:

$$V_i(A_i) \geq \frac{\sqrt{n}}{\sqrt{n} + |S|} V_i(A_i^*) \geq \frac{\sqrt{n}}{n} V_i(A_i^*) = \frac{1}{\sqrt{n}} V_i(A_i^*).$$

- For $i \in S$, $V_i(A_i^*) < \frac{1}{\sqrt{n}}$ (by definition), and $V_i(A_i) \geq \frac{1}{n}$ (by proportionality we just built). We have

$$\frac{V_i(A_i^*)}{\sqrt{n}} < \frac{1}{n} \leq V_i(A_i).$$



Case 1: Why the loss is at most $\sqrt{n}$

Moreover,

- for $i \in L$:

$$V_i(A_i) \geq \frac{\sqrt{n}}{\sqrt{n} + |S|} V_i(A_i^*) \geq \frac{\sqrt{n}}{n} V_i(A_i^*) = \frac{1}{\sqrt{n}} V_i(A_i^*).$$

- For $i \in S$, $V_i(A_i^*) < \frac{1}{\sqrt{n}}$ (by definition), and $V_i(A_i) \geq \frac{1}{n}$ (by proportionality we just built). We have

$$\frac{V_i(A_i^*)}{\sqrt{n}} < \frac{1}{n} \leq V_i(A_i).$$


Therefore,

$$\text{sw}(A) = \sum_i V_i(A_i) \geq \frac{1}{\sqrt{n}} \sum_{i \in N} V_i(A_i^*) = \frac{\text{sw}(A^*)}{\sqrt{n}}.$$



Lower Bound Idea for Proportionality


- Say that there are $\sqrt{n}$ *large* agents. Each agent $i \in L$ values **only one** distinct interval L_i of length $1/\sqrt{n}$ uniformly.



- The remaining agents (S) value the entire cake uniformly.
- Unconstrained welfare serves the large agents $\Rightarrow \text{sw}(\mathbf{A}^*) = \sqrt{n}$.
- Any proportional $\mathbf{A}$ gives an interval of length $1/n$ to each *small agent*, leaving only $(1 - (n - \sqrt{n})/n) = 1/\sqrt{n}$ to agents in L .

Lower Bound Idea for Proportionality

- Say that there are $\sqrt{n}$ *large* agents. Each agent $i \in L$ values **only one** distinct interval L_i of length $1/\sqrt{n}$ uniformly.



- The remaining agents (S) value the entire cake uniformly.
- Unconstrained welfare serves the large agents $\Rightarrow \text{sw}(\mathbf{A}^*) = \sqrt{n}$.
- Any proportional $\mathbf{A}$ gives an interval of length $1/n$ to each *small agent*, leaving only $(1 - (n - \sqrt{n})/n) = 1/\sqrt{n}$ to agents in L .

$$\sum_{i \in L} V_i(A_i) + \sum_{j \in S} V_j(A_j) \leq \sqrt{n}/\sqrt{n} + 1 \leq 2.$$



Price of Envy-Freeness and Equitability

- Since envy-freeness implies proportionality, the $\Omega(\sqrt{n})$ lower bound also applies to envy-freeness.
- No nontrivial $o(n)$ upper bound on the price of envy-freeness is known.



Price of Envy-Freeness and Equitability

- Since envy-freeness implies proportionality, the $\Omega(\sqrt{n})$ lower bound also applies to envy-freeness.
- No nontrivial $o(n)$ upper bound on the price of envy-freeness is known.
- Equitability can be much more expensive.

Theorem (Caragiannis et al., WINE 2009)

The price of equitability is $\Theta(n)$.



Equitable vs. Envy-Free Welfare

Theorem (Brams et al., AAI 2012)

For piecewise constant valuations, the optimal envy-free allocation has social welfare at least that of the optimal equitable allocation.

Optimal equitable
allocation

$$sw(EF) \geq sw(EQ)$$

Optimal envy-free
allocation



Proof Tool: Linear Fisher Markets

A linear Fisher market has:

- divisible goods G ;
- utilities u_{ij} for agent i and good j ;
- budgets e_i ;
- equilibrium prices p_j
- allocations f_{ij} (giving agent i a fraction f_{ij} of good j).
- utility of agent i for an allocation: $\sum_{j \in G} f_{ij} u_{ij}$.
- For all $j \in G$, $\sum_{i \in N} f_{ij} \leq 1$.



Fisher-Market View of Cake Cutting

Start from an **equitable** welfare-maximizing cake allocation

$$\mathbf{A}^* = (A_1^*, A_2^*, \dots, A_n^*).$$

Construct a linear Fisher market

- Good j corresponds to cake piece A_j^* .
- Agent i 's utility for good j is $u_{ij} = V_i(A_j^*)$.
- Budgets are identical: $e_i = \frac{1}{n} \quad \forall i \in N$.

cake pieces A_j^* $\longleftrightarrow$ divisible Fisher goods j .



From Fisher Allocation Back to Cake

- Given a Fisher allocation $\mathbf{f} = (f_{ij})_{i \in N, j \in G}$, we want a cake allocation $\mathbf{A}$ such that

$$V_i(A_j) = \sum_{k \in G} f_{jk} u_{ik} \quad \forall i, j \in N.$$

Idea

Subdivide each A_k^* so that, within every resulting interval, each agent's value is proportional to length.

- The utilities in the Fisher market can be replicated in the cake-cutting instance.



Why the Fisher Allocation Is Envy-Free?

- Let $p = (p_1, p_2, \dots, p_n)$ be **equilibrium prices**, and define $r_i^* = \max_{j \in G} \frac{u_{ij}}{p_j}$.
- By the Fisher-market property (Vazirani 2007), agent i receives only **goods maximizing utility per price** (bang-per-buck):

$$f_{ij} > 0 \quad \implies \quad \frac{u_{ij}}{p_j} = r_i^*.$$



Why the Fisher Allocation Is Envy-Free?

- Let $p = (p_1, p_2, \dots, p_n)$ be **equilibrium prices**, and define $r_i^* = \max_{j \in G} \frac{u_{ij}}{p_j}$.
- By the Fisher-market property (Vazirani 2007), agent i receives only **goods maximizing utility per price** (bang-per-buck):

$$f_{ij} > 0 \implies \frac{u_{ij}}{p_j} = r_i^*.$$

- Hence (note: $\sum_{j \in G} f_{ij} p_j = e_i := \frac{1}{n}$; agents spend their entire budgets)

$$V_i(A_i) = \sum_{j \in G} f_{ij} u_{ij} = \sum_{j \in G} f_{ij} \frac{u_{ij}}{p_j} p_j = \sum_{j \in G} f_{ij} r_i^* p_j = \frac{r_i^*}{n}.$$



Why the Fisher Allocation Is Envy-Free?

- Let $p = (p_1, p_2, \dots, p_n)$ be **equilibrium prices**, and define $r_i^* = \max_{j \in G} \frac{u_{ij}}{p_j}$.
- By the Fisher-market property (Vazirani 2007), agent i receives only **goods maximizing utility per price** (bang-per-buck):

$$f_{ij} > 0 \implies \frac{u_{ij}}{p_j} = r_i^*.$$

- Hence (note: $\sum_{j \in G} f_{ij} p_j = e_i := \frac{1}{n}$; agents spend their entire budgets)

$$V_i(A_i) = \sum_{j \in G} f_{ij} u_{ij} = \sum_{j \in G} f_{ij} \frac{u_{ij}}{p_j} p_j = \sum_{j \in G} f_{ij} r_i^* p_j = \frac{r_i^*}{n}.$$

- For another agent k ,

$$V_i(A_k) = \sum_{j \in G} f_{kj} u_{ij} \leq \sum_{j \in G} f_{kj} r_i^* p_j = \frac{r_i^*}{n}.$$



Why the Fisher Allocation Is Envy-Free?

- Let $p = (p_1, p_2, \dots, p_n)$ be **equilibrium prices**, and define $r_i^* = \max_{j \in G} \frac{u_{ij}}{p_j}$.
- By the Fisher-market property (Vazirani 2007), agent i receives only **goods maximizing utility per price** (bang-per-buck):

$$f_{ij} > 0 \implies \frac{u_{ij}}{p_j} = r_i^*.$$

- Hence (note: $\sum_{j \in G} f_{ij} p_j = e_i := \frac{1}{n}$; agents spend their entire budgets)

$$V_i(A_i) = \sum_{j \in G} f_{ij} u_{ij} = \sum_{j \in G} f_{ij} \frac{u_{ij}}{p_j} p_j = \sum_{j \in G} f_{ij} r_i^* p_j = \frac{r_i^*}{n}.$$

- For another agent k ,

$$V_i(A_k) = \sum_{j \in G} f_{kj} u_{ij} \leq \sum_{j \in G} f_{kj} r_i^* p_j = \frac{r_i^*}{n}.$$

- Therefore $V_i(A_i) \geq V_i(A_k)$, so $\mathbf{A}$ is envy-free.



Why Welfare Does Not Decrease

- Since $\mathbf{A}^*$ is equitable, there exists $\alpha > 0$ such that $V_i(A_i^*) = \alpha \ \forall i \in N$.

Thus $\text{sw}(\mathbf{A}^*) = n\alpha$.



Why Welfare Does Not Decrease

- Since $\mathbf{A}^*$ is equitable, there exists $\alpha > 0$ such that $V_i(A_i^*) = \alpha \ \forall i \in N$.

Thus $\text{sw}(\mathbf{A}^*) = n\alpha$.

- From the previous page, $\text{sw}(\mathbf{A}) = \frac{1}{n} \sum_{i \in N} r_i^*$. (recall: $\mathbf{A}$ is the Fisher allocation)



Why Welfare Does Not Decrease

- Since $\mathbf{A}^*$ is equitable, there exists $\alpha > 0$ such that $V_i(A_i^*) = \alpha \ \forall i \in N$.

Thus $\text{sw}(\mathbf{A}^*) = n\alpha$.

- From the previous page, $\text{sw}(\mathbf{A}) = \frac{1}{n} \sum_{i \in N} r_i^*$. (recall: $\mathbf{A}$ is the Fisher allocation)

- Since $r_i^* = \max_{j \in G} \frac{u_{ij}}{p_j} \geq \frac{u_{ii}}{p_i} = \frac{\alpha}{p_i}$, we get



Why Welfare Does Not Decrease

- Since $\mathbf{A}^*$ is equitable, there exists $\alpha > 0$ such that $V_i(A_i^*) = \alpha \quad \forall i \in N$.

Thus $\text{sw}(\mathbf{A}^*) = n\alpha$.

- From the previous page, $\text{sw}(\mathbf{A}) = \frac{1}{n} \sum_{i \in N} r_i^*$. (recall: $\mathbf{A}$ is the Fisher allocation)

- Since $r_i^* = \max_{j \in G} \frac{u_{ij}}{p_j} \geq \frac{u_{ii}}{p_i} = \frac{\alpha}{p_i}$, we get $\text{sw}(\mathbf{A}) \geq \frac{1}{n} \sum_{i \in N} \frac{\alpha}{p_i} = \frac{\alpha}{n} \sum_{i \in N} \frac{1}{p_i}$.



Why Welfare Does Not Decrease

- Since $\mathbf{A}^*$ is equitable, there exists $\alpha > 0$ such that $V_i(A_i^*) = \alpha \quad \forall i \in N$.

Thus $\text{sw}(\mathbf{A}^*) = n\alpha$.

- From the previous page, $\text{sw}(\mathbf{A}) = \frac{1}{n} \sum_{i \in N} r_i^*$. (recall: $\mathbf{A}$ is the Fisher allocation)

- Since $r_i^* = \max_{j \in G} \frac{u_{ij}}{p_j} \geq \frac{u_{ii}}{p_i} = \frac{\alpha}{p_i}$, we get $\text{sw}(\mathbf{A}) \geq \frac{1}{n} \sum_{i \in N} \frac{\alpha}{p_i} = \frac{\alpha}{n} \sum_{i \in N} \frac{1}{p_i}$.

$$\sum_{i \in N} p_i = 1, \Rightarrow \sum_{i \in N} \frac{1}{p_i} \geq n^2 \text{ (Cauchy-Schwarz (Exercise)).}$$

- Therefore, $\text{sw}(\mathbf{A}) \geq \frac{\alpha}{n} \cdot n^2 = n\alpha = \text{sw}(\mathbf{A}^*)$.



Takeaways

- 1 Proportionality is algorithmically well understood in the Robertson–Webb model: $\Theta(n \log n)$ queries.
- 2 Envy-freeness is strictly harder; exact bounded algorithms are much more delicate.
- 3 Explicit piecewise constant valuations enable LP-based computation of optimal fair allocations.
- 4 Fairness can significantly reduce social welfare; the price depends strongly on the fairness notion.



Discussions

