Mathematics for Machine Learning
— Continuous Optimization
Gradient Descent and Constrained Optimization

Joseph Chuang-Chieh Lin

Department of Computer Science & Engineering,
National Taiwan Ocean University

Fall 2025

Credits for the resource

- The slides are based on the textbooks:

  - *Marc Peter Deisenroth, A. Aldo Faisal, and Cheng Soon Ong: Mathematics for Machine Learning. Cambridge University Press. 2020.*
  - *Arnold J. Insel, Lawrence E. Spence, Stephen H. Friedberg: Linear Algebra, 4th Edition. Prentice Hall. 2013.*
  - *Howard Anton, Chris Rorres, Anton Kaul: Elementary Linear Algebra, 12th Edition. Wiley. 2019.*

- We could partially refer to the monograph:
  Francesco Orabona: A Modern Introduction to Online Learning.
  https://arxiv.org/abs/1912.13213

Outline

Motivation

- Machine learning algorithms are solving mathematical formulations which are expressed as numerical optimization methods.

- We focus on basic numerical methods for training machine learning models.
  - This boils down to *finding a "good" set of parameters*.
  - Goodness: determined by the objective function or the probabilistic model.

- Given an objective function, finding the best value of parameters is done using optimization algorithms.

Remark

- We will discuss two branches of continuous optimization:
  - Unconstrained optimization.
  - Constrained optimization.

- Assume that the objective functions are differentiable.

- We focus on "minimization" objectives.

- We will make use of the "gradients".

Example

Consider the loss function $\ell(x) = x^4 + 7x^3 + 5x^2 - 17x + 3$.

The gradient:
$$\frac{\mathrm{d}\ell(x)}{\mathrm{d}x} = 4x^3 + 21x^2 + 10x - 17.$$

The second derivative:
$$\frac{\mathrm{d}^2\ell(x)}{\mathrm{d}x^2} = 12x^2 + 42x + 10.$$

Solving $\frac{\mathrm{d}\ell(x)}{\mathrm{d}x} = 0$ we get $x = -4.5, -1.4$, or 0.7.

Example

Consider the loss function $\ell(x) = x^4 + 7x^3 + 5x^2 - 17x + 3$.

The gradient:
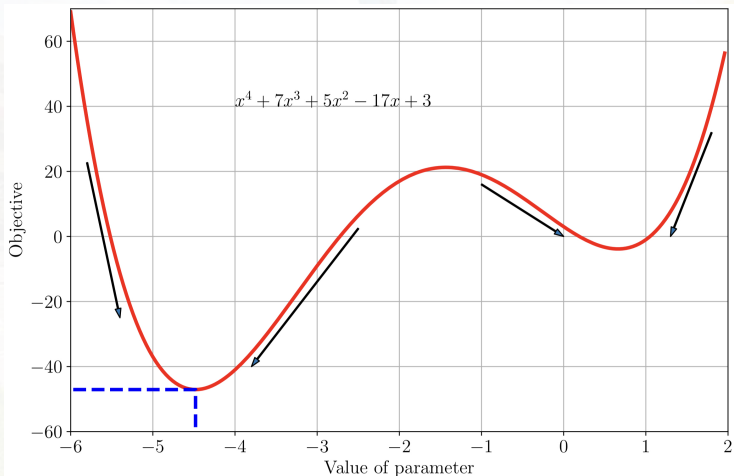$$\frac{\mathrm{d}\ell(x)}{\mathrm{d}x} = 4x^3 + 21x^2 + 10x - 17.$$

The second derivative:
$$\frac{\mathrm{d}^2\ell(x)}{\mathrm{d}x^2} = 12x^2 + 42x + 10.$$

Solving $\frac{\mathrm{d}\ell(x)}{\mathrm{d}x} = 0$ we get $x = -4.5, -1.4,$ or 0.7.

By checking whether $\frac{\mathrm{d}^2\ell(x)}{\mathrm{d}x^2}$ is positive or negative at the stationary point(s), we know $x = -1.4$ is a (local) maximum.

Function Plot & Negative Gradients of Univariate $\ell(x)$

Start at some x_0, and then the negative gradient leads us to some (local) minimum.

$$x^4 + 7x^3 + 5x^2 - 17x + 3$$

Heads Up

- For convex functions, there is no such a tricky dependency on the starting point.
  - All local minimums are *global* minimum.

Heads Up

- For convex functions, there is no such a tricky dependency on the starting point.
    - All local minimums are *global* minimum.

- For "maximization" objectives, we shall follow the (positive) gradients.
    - Minimization objective $\implies$ follows the negative gradient $\implies$ "gradient descent".
    - Maximization objective $\implies$ follows the (positive) gradient $\implies$ "gradient ascent".

Heads Up

- For convex functions, there is no such a tricky dependency on the starting point.
    - All local minimums are *global* minimum.

- For "maximization" objectives, we shall follow the (positive) gradients.
    - Minimization objective $\implies$ follows the negative gradient $\implies$ "gradient descent".
    - Maximization objective $\implies$ follows the (positive) gradient $\implies$ "gradient ascent".

- For optimization in higher dimensions, it is almost impossible to visualize the idea of gradients, descent directions and optimal values.

Outline

The Problem

> **Solving for the minimum of a real-valued function**
>
> $$\min_{\mathbf{x}} f(\mathbf{x}),$$
>
> where $f \colon \mathbb{R}^d \to \mathbb{R}$ is the objective function which is assumed to be differentiable.

Gradient Descent

- Gradient descent is a first-order optimization algorithm.

Gradient Descent

- Starting at a particular location $\mathbf{x}_0$.

The algorithm runs iteratively by giving

$$\mathbf{x}_{t+1} = \mathbf{x}_t - \gamma_t((\nabla f)(\mathbf{x}_t)).$$

where $\gamma \geq 0$ is called the step-size (or learning rate).

Goal: $f(\mathbf{x}_0) \geq f(\mathbf{x}_1) \geq \cdots$ converges to a local minimum.

Example (1/2)

Example

Consider

$$f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) = \frac{1}{2} \left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]^\top \left[\begin{array}{cc} 2 & 1 \\ 1 & 20 \end{array}\right] \left[\begin{array}{c} x_1 \\ x_2 \end{array}\right] - \left[\begin{array}{c} 5 \\ 3 \end{array}\right]^\top \left[\begin{array}{c} x_1 \\ x_2 \end{array}\right].$$

Compute $\nabla f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) =$

Example (1/2)

Example

Consider

$$f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) = \frac{1}{2}\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]^{\top}\left[\begin{array}{cc} 2 & 1 \\ 1 & 20 \end{array}\right]\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right] - \left[\begin{array}{c} 5 \\ 3 \end{array}\right]^{\top}\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right].$$

Compute $\nabla f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) =$

Example (1/2)

Example

Consider

$$f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) = \frac{1}{2}\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]^{\top}\left[\begin{array}{cc} 2 & 1 \\ 1 & 20 \end{array}\right]\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right] - \left[\begin{array}{c} 5 \\ 3 \end{array}\right]^{\top}\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right].$$

Compute $\nabla f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) = \left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]^{\top}\left[\begin{array}{cc} 2 & 1 \\ 1 & 20 \end{array}\right] - \left[\begin{array}{c} 5 \\ 3 \end{array}\right]^{\top}.$

Example (1/2)

Example

Consider

$$f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) = \frac{1}{2}\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]^\top \left[\begin{array}{cc} 2 & 1 \\ 1 & 20 \end{array}\right]\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right] - \left[\begin{array}{c} 5 \\ 3 \end{array}\right]^\top \left[\begin{array}{c} x_1 \\ x_2 \end{array}\right].$$

Compute $\nabla f\left(\left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]\right) = \left[\begin{array}{c} x_1 \\ x_2 \end{array}\right]^\top \left[\begin{array}{cc} 2 & 1 \\ 1 & 20 \end{array}\right] - \left[\begin{array}{c} 5 \\ 3 \end{array}\right]^\top.$

Running gradient descent and starting at $\mathbf{x}_0 = [-3, \ -1]^\top$, what's $\mathbf{x}_1$?
And what's $\mathbf{x}_2$?

Example (2/2)

Remark

The gradient points in a direction orthogonal to the contour lines of the function.

Remark

The gradient points in a direction orthogonal to the contour lines of the function.

- Let $\gamma(t) = (x(t), y(y)) \in \mathbb{R}^2$ be a curve.
  - A contour of f: $f(\gamma(t)) = C$ for some constant $C \in \mathbb{R}$.

Remark

The gradient points in a direction orthogonal to the contour lines of the function.

- Let $\gamma(t) = (x(t), y(y)) \in \mathbb{R}^2$ be a curve.
  - A contour of f: $f(\gamma(t)) = C$ for some constant $C \in \mathbb{R}$.

$$\frac{\mathrm{d}f}{\mathrm{d}t} = \mathbf{0}.$$

But

$$\frac{\mathrm{d}f}{\mathrm{d}t} \quad = \quad \frac{\mathrm{d}f}{\mathrm{d}\gamma}\frac{\mathrm{d}\gamma}{\mathrm{d}t} = \begin{bmatrix} \frac{\partial f}{\partial x} & \frac{\partial f}{\partial y} \end{bmatrix} \begin{bmatrix} \frac{\partial x}{\partial t} \\ \frac{\partial y}{\partial t} \end{bmatrix} = \langle \nabla_\gamma f, \nabla_t \gamma(t) \rangle$$

On the Step-size

- Adaptive gradient descent: rescale the step-size γ at each iteration.

On the Step-size

- Adaptive gradient descent: rescale the step-size γ at each iteration.

- Two simple heuristics:
  - When the function value $\uparrow$ after a gradient step $\Rightarrow$ undo the step and decrease the step-size.
  - When the function value $\downarrow$ after a gradient step $\Rightarrow$ try to increase the step-size.

Another Example: Solving a Linear Equation System

Example

Consider $\boldsymbol{A}\mathbf{x} = \mathbf{b}$.

Another Example: Solving a Linear Equation System

Example

Consider $\boldsymbol{A}\mathbf{x} = \mathbf{b}$. $\implies$ $\boldsymbol{A}\mathbf{x} - \mathbf{b} = \mathbf{0}$.

- We want to find the solution *approximately* by minimizing the squared error

$$\ell(\mathbf{x}) = \|\boldsymbol{A}\mathbf{x} - \mathbf{b}\|^2 = (\boldsymbol{A}\mathbf{x} - \mathbf{b})^\top (\boldsymbol{A}\mathbf{x} - \mathbf{b})$$

  where $\|\cdot\|$ is the ℓ_2-norm.

- $\nabla_{\mathbf{x}}\ell(\mathbf{x}) =$

Another Example: Solving a Linear Equation System

Example

Consider $\mathbf{A}\mathbf{x} = \mathbf{b}$. $\implies$ $\mathbf{A}\mathbf{x} - \mathbf{b} = \mathbf{0}$.

- We want to find the solution *approximately* by minimizing the equared error

$$\ell(\mathbf{x}) = \|\mathbf{A}\mathbf{x} - \mathbf{b}\|^2 = (\mathbf{A}\mathbf{x} - \mathbf{b})^\top (\mathbf{A}\mathbf{x} - \mathbf{b})$$

where $\|\cdot\|$ is the ℓ_2-norm.

- $\nabla_{\mathbf{x}}\ell(\mathbf{x}) = 2(\mathbf{A}\mathbf{x} - \mathbf{b})^\top \mathbf{A}$, run the gradient descent algorithm.

Another Example: Solving a Linear Equation System

Example

Consider $\boldsymbol{A}\mathbf{x} = \mathbf{b}$. $\implies$ $\boldsymbol{A}\mathbf{x} - \mathbf{b} = \mathbf{0}$.

- We want to find the solution *approximately* by minimizing the equared error

$$\ell(\mathbf{x}) = \|\boldsymbol{A}\mathbf{x} - \mathbf{b}\|^2 = (\boldsymbol{A}\mathbf{x} - \mathbf{b})^\top (\boldsymbol{A}\mathbf{x} - \mathbf{b})$$

  where $\|\cdot\|$ is the ℓ_2-norm.

- $\nabla_\mathbf{x}\ell(\mathbf{x}) = 2(\boldsymbol{A}\mathbf{x} - \mathbf{b})^\top \boldsymbol{A}$, run the gradient descent algorithm.

Okay, this is just an example and somehow not a good idea because there is an analytic solution.

Another Example: Solving a Linear Equation System

Example

Consider $\boldsymbol{A}\mathbf{x} = \mathbf{b}. \implies \boldsymbol{A}\mathbf{x} - \mathbf{b} = \mathbf{0}$.

- We want to find the solution *approximately* by minimizing the equared error

$$\ell(\mathbf{x}) = \|\boldsymbol{A}\mathbf{x} - \mathbf{b}\|^2 = (\boldsymbol{A}\mathbf{x} - \mathbf{b})^{\top}(\boldsymbol{A}\mathbf{x} - \mathbf{b})$$

  where $\| \cdot \|$ is the ℓ_2-norm.

- $\nabla_{\mathbf{x}}\ell(\mathbf{x}) = 2(\boldsymbol{A}\mathbf{x} - \mathbf{b})^{\top}\boldsymbol{A}$, run the gradient descent algorithm.

Okay, this is just an example and somehow not a good idea because there is an analytic solution. (Solving $\nabla_{\mathbf{x}}\ell(\mathbf{x}) = \mathbf{0}$)

Outline

1. Preface / Introduction

2. Optimization Using Gradient Descent
   - Gradient Descent with Momentum
   - Stochastic Gradient Descent

3. Constrained Optimization

Gradient Descent with Momentum

- The convergence of gradient descent could be slow due to the curvature of the optimization surface.

- **Idea:** Give gradient descent some *memory*.
  - Introducing an additional term to remember what happened in the previous iteration.

- The steps (for $\alpha \in [0, 1]$):

$$
\begin{aligned}
\mathbf{x}_{t+1} &= \mathbf{x}_t - \gamma_t((\nabla f)(\mathbf{x}_t))^\top + \alpha \Delta \mathbf{x}_t \\
\Delta \mathbf{x}_t &= \mathbf{x}_t - \mathbf{x}_{t-1}
\end{aligned}
$$

Gradient Descent with Momentum

- The convergence of gradient descent could be slow due to the curvature of the optimization surface.

- **Idea:** Give gradient descent some *memory*.
  - Introducing an additional term to remember what happened in the previous iteration.

- The steps (for $\alpha \in [0, 1]$):

$$
\begin{aligned}
\mathbf{x}_{t+1} &= \mathbf{x}_t - \gamma_t((\nabla f)(\mathbf{x}_t))^\top + \alpha \Delta \mathbf{x}_t \\
\Delta \mathbf{x}_t &= \mathbf{x}_t - \mathbf{x}_{t-1} = \alpha \Delta \mathbf{x}_{t-1} - \gamma_{t-1}((\nabla f)(\mathbf{x}_{t-1}))^\top
\end{aligned}
$$

Outline

1. Preface / Introduction

2. Optimization Using Gradient Descent
   - Gradient Descent with Momentum
   - **Stochastic Gradient Descent**

3. Constrained Optimization

Stochastic Gradient Descent (1/5)

Motivation:

- Computing the gradient can be very time consuming.
- Approximating the gradient is useful.
  - We aim at only knowing a noisy approximation to the gradient.

Stochastic Gradient Descent (2/5)

The objective function:

$$L(\boldsymbol{\theta}) = \sum_{i=1}^{N} L_i(\boldsymbol{\theta}),$$

which is sum of losses L_i incurred by each sample i. $\boldsymbol{\theta}$ is the vector of parameters of interest.

- **Goal:** Find $\boldsymbol{\theta}$ that minimizes L.

Example: log-likelihoods

$$L(\boldsymbol{\theta}) = -\sum_{i=1}^{N} \log p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}),$$

for the training inputs $\mathbf{x}_i \in \mathbb{R}^D$, training targets y_i, and the parameters $\boldsymbol{\theta}$ of the model.

Stochastic Gradient Descent (3/5)

Updating $\boldsymbol{\theta}$:

$$\boldsymbol{\theta}_{t+1} \leftarrow \boldsymbol{\theta}_t - \gamma_t (\nabla L(\boldsymbol{\theta}_t))^\top$$

Stochastic Gradient Descent (3/5)

Updating $\boldsymbol{\theta}$:

$$\boldsymbol{\theta}_{t+1} \leftarrow \boldsymbol{\theta}_t - \gamma_t (\nabla L(\boldsymbol{\theta}_t))^\top = \boldsymbol{\theta}_t - \gamma_t \sum_{i=1}^{N} (\nabla L_i(\boldsymbol{\theta}_t))^\top$$

according to suitable γ_t's.

Stochastic Gradient Descent (3/5)

Updating $\boldsymbol{\theta}$:   (Note that the differential operator is linear)

$$\boldsymbol{\theta}_{t+1} \leftarrow \boldsymbol{\theta}_t - \gamma_t (\nabla L(\boldsymbol{\theta}_t))^\top = \boldsymbol{\theta}_t - \gamma_t \sum_{i=1}^{N} (\nabla L_i(\boldsymbol{\theta}_t))^\top$$

according to suitable γ_t's.

Issues

When training set is enormous or no simple formulas exist for evaluating the (sum of) gradients.

Stochastic Gradient Descent (4/5)

Idea: Consider taking a sum of a smaller set of L_n.

- For $i = 1, 2, \ldots, N$, randomly choose a subset of L_i for mini-batch gradient descent.

Stochastic Gradient Descent (4/5)

Idea: Consider taking a sum of a smaller set of L_n.

- For $i = 1, 2, \ldots, N$, randomly choose a subset of L_i for mini-batch gradient descent.

- Reasons:
  - We only require the gradient to be an unbiased estimate of the true gradient.

Stochastic Gradient Descent (4/5)

Idea: Consider taking a sum of a smaller set of L_n.

- For $i = 1, 2, \ldots, N$, randomly choose a subset of L_i for mini-batch gradient descent.

- Reasons:
  - We only require the gradient to be an unbiased estimate of the true gradient.
  - Even the term $\sum_{i=1}^{N}(\nabla L_i(\boldsymbol{\theta}_i))$ is just an empirical estimate of the expected value of the gradient.

Stochastic Gradient Descent (4/5)

Idea: Consider taking a sum of a smaller set of L_n.

- For $i = 1, 2, \ldots, N$, randomly choose a subset of L_i for mini-batch gradient descent.

- Reasons:
    - We only require the gradient to be an unbiased estimate of the true gradient.
    - Even the term $\sum_{i=1}^{N}(\nabla L_i(\boldsymbol{\theta}_i))$ is just an empirical estimate of the expected value of the gradient.

- Benefits for mini-batch:

Stochastic Gradient Descent (4/5)

Idea: Consider taking a sum of a smaller set of L_n.

- For $i = 1, 2, \ldots, N$, randomly choose a subset of L_i for mini-batch gradient descent.

- Reasons:
    - We only require the gradient to be an unbiased estimate of the true gradient.
    - Even the term $\sum_{i=1}^{N}(\nabla L_i(\boldsymbol{\theta}_i))$ is just an empirical estimate of the expected value of the gradient.

- Benefits for mini-batch:
    - Quick to estimate.
    - Noisy estimate allows us to get out of some bad local optima.
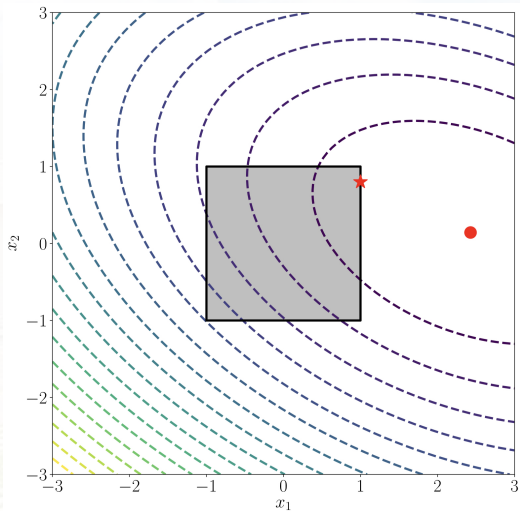
Stochastic Gradient Descent (4/5)

Idea: Consider taking a sum of a smaller set of L_n.

- For $i = 1, 2, \ldots, N$, randomly choose a subset of L_i for mini-batch gradient descent.

- Reasons:
  - We only require the gradient to be an unbiased estimate of the true gradient.
  - Even the term $\sum_{i=1}^{N}(\nabla L_i(\boldsymbol{\theta}_i))$ is just an empirical estimate of the expected value of the gradient.

- Benefits for mini-batch:
  - Quick to estimate.
  - Noisy estimate allows us to get out of some bad local optima.
  - Good for generalization.

Outline

Constrained Optimization

Constrained Optimization

The objective function:

$$\min_{\mathbf{x}} \quad f(\mathbf{x})$$
$$\text{subject to} \quad g_i(\mathbf{x}) \leq 0, \quad \text{for all } i = 1, \ldots, m.$$

Note: f and g_i could be non-convex in general.

Constrained Optimization

The objective function:

$$\min_{\mathbf{x}} \quad f(\mathbf{x})$$
$$\text{subject to} \quad g_i(\mathbf{x}) \leq 0, \quad \text{for all } i = 1, \ldots, m.$$

Note: f and g_i could be non-convex in general.

An Easy Unconstrained Objective

$$J(\mathbf{x}) = f(\mathbf{x}) + \sum_{i=1}^{m} \mathbf{1}(g_i(\mathbf{x})),$$

where $\mathbf{1}(z)$ is an infinite step function $\mathbf{1}(z) = \begin{cases} 0 & \text{if } z \leq 0 \\ \infty & \text{otherwise} \end{cases}$.

Constrained Optimization

The objective function:

$$\min_{\mathbf{x}} \quad f(\mathbf{x})$$
$$\text{subject to} \quad g_i(\mathbf{x}) \leq 0, \quad \text{for all } i = 1, \ldots, m.$$

Note: f and g_i could be non-convex in general.

An Easy Unconstrained Objective

$$J(\mathbf{x}) = f(\mathbf{x}) + \sum_{i=1}^{m} \mathbf{1}(g_i(\mathbf{x})),$$

where $\mathbf{1}(z)$ is an infinite step function $\mathbf{1}(z) = \begin{cases} 0 & \text{if } z \leq 0 \\ \infty & \text{otherwise} \end{cases}$.

The infinite step function is difficult to optimize...

A Workaround Approach: Lagrange Multipliers

For $\boldsymbol{\lambda}_i \geq 0$, define

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) \;=\; f(\mathbf{x}) + \sum_{i=1}^{m} \lambda_i g_i(\mathbf{x})$$

A Workaround Approach: Lagrange Multipliers

For $\boldsymbol{\lambda}_i \geq 0$, define

$$
\begin{aligned}
\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) &= f(\mathbf{x}) + \sum_{i=1}^{m} \lambda_i g_i(\mathbf{x}) \\
&= f(\mathbf{x}) + \boldsymbol{\lambda}^{\top} \boldsymbol{g}(\mathbf{x}).
\end{aligned}
$$

A Workaround Approach: Lagrange Multipliers

For $\lambda_i \geq 0$, define

$$\begin{aligned}
\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) &= f(\mathbf{x}) + \sum_{i=1}^{m} \lambda_i g_i(\mathbf{x}) \\
&= f(\mathbf{x}) + \boldsymbol{\lambda}^{\top} \mathbf{g}(\mathbf{x}).
\end{aligned}$$

- We concatenate all constraints $g_i(\mathbf{x})$ into a vector $\mathbf{g}(\mathbf{x})$.

A Workaround Approach: Lagrange Multipliers

For $\boldsymbol{\lambda}_i \geq 0$, define

$$
\begin{aligned}
\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) &= f(\mathbf{x}) + \sum_{i=1}^{m} \lambda_i g_i(\mathbf{x}) \\
&= f(\mathbf{x}) + \boldsymbol{\lambda}^\top \boldsymbol{g}(\mathbf{x}).
\end{aligned}
$$

- We concatenate all constraints $g_i(\mathbf{x})$ into a vector $\boldsymbol{g}(\mathbf{x})$.
  - This technique is widely used in building machine learning models!

A Workaround Approach: Lagrange Multipliers

For $\boldsymbol{\lambda}_i \geq 0$, define

$$
\begin{aligned}
\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) &= f(\mathbf{x}) + \sum_{i=1}^{m} \lambda_i g_i(\mathbf{x}) \\
&= f(\mathbf{x}) + \boldsymbol{\lambda}^{\top} \boldsymbol{g}(\mathbf{x}).
\end{aligned}
$$

- We concatenate all constraints $g_i(\mathbf{x})$ into a vector $\boldsymbol{g}(\mathbf{x})$.
    - This technique is widely used in building machine learning models!

> - $\mathbf{x}$: primal variables.
> - $\boldsymbol{\lambda}$: dual variables.

Primal & Dual Problems

The primal problem

$$\min_{\mathbf{x}} \quad f(\mathbf{x})$$
subject to $\quad g_i(\mathbf{x}) \leq 0, \text{ for } i \in [m].$

The dual problem

$$\max_{\boldsymbol{\lambda} \in \mathbb{R}^m} \quad \mathcal{D}(\boldsymbol{\lambda})$$
subject to $\quad \boldsymbol{\lambda} \geq \mathbf{0}.$

$\mathcal{D}(\boldsymbol{\lambda}) := \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$

Primal & Dual Problems

The primal problem
$$\min_{\mathbf{x}} \quad f(\mathbf{x})$$ subject to $\quad g_i(\mathbf{x}) \leq 0$, for $i \in [m]$.

The dual problem
$$\max_{\boldsymbol{\lambda} \in \mathbb{R}^m} \quad \mathcal{D}(\boldsymbol{\lambda})$$ subject to $\quad \boldsymbol{\lambda} \geq \mathbf{0}$.

$\mathcal{D}(\boldsymbol{\lambda}) := \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$.

- This is unconstrained optimization problem for a given value of $\boldsymbol{\lambda}$.

  ($\boldsymbol{\lambda} \geq \mathbf{0} \Leftrightarrow \lambda_i \geq 0$ for each i

Primal & Dual Problems

The primal problem

$$\min_{\mathbf{x}} \quad f(\mathbf{x})$$
$$\text{subject to} \quad g_i(\mathbf{x}) \leq 0, \text{ for } i \in [m].$$

The dual problem

$$\max_{\boldsymbol{\lambda} \in \mathbb{R}^m} \quad \mathcal{D}(\boldsymbol{\lambda})$$
$$\text{subject to} \quad \boldsymbol{\lambda} \geq \mathbf{0}.$$

$\mathcal{D}(\boldsymbol{\lambda}) := \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$

- This is unconstrained optimization problem for a given value of $\boldsymbol{\lambda}$.

  ($\boldsymbol{\lambda} \geq \mathbf{0} \Leftrightarrow \lambda_i \geq 0$ for each $i \Leftrightarrow \boldsymbol{\lambda} \succeq \mathbf{0}$)

Minimax Inequality

Minimax Inequality

For any function φ with two arguments $\mathbf{x}, \mathbf{y}$,

$$\max_{\mathbf{y}} \min_{\mathbf{x}} \varphi(\mathbf{x}, \mathbf{y}) \leq \min_{\mathbf{x}} \max_{\mathbf{y}} \varphi(\mathbf{x}, \mathbf{y}).$$

Minimax Inequality

Minimax Inequality

For any function φ with two arguments $\mathbf{x}, \mathbf{y}$,

$$\max_{\mathbf{y}} \min_{\mathbf{x}} \varphi(\mathbf{x}, \mathbf{y}) \leq \min_{\mathbf{x}} \max_{\mathbf{y}} \varphi(\mathbf{x}, \mathbf{y}).$$

- Consider the inequality

$$\text{For all } \mathbf{x}_0, \mathbf{y}_0, \quad \min_{\mathbf{x}} \varphi(\mathbf{x}, \mathbf{y}_0) \leq \max_{\mathbf{y}} \varphi(\mathbf{x}_0, \mathbf{y}).$$

- This implies that

$$\min_{\mathbf{x} \in \mathbb{R}^d} \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) \geq \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$$

Compare $J(\mathbf{x})$ with $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$

$$J(\mathbf{x}) = f(\mathbf{x}) + \sum_{i=1}^{m} \mathbf{1}(g_i(\mathbf{x})), \text{ where } \mathbf{1}(z) = \begin{cases} 0 & \text{if } z \leq 0 \\ \infty & \text{otherwise} \end{cases}$$

v.s.

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \boldsymbol{\lambda}^\top \boldsymbol{g}(\mathbf{x}).$$

Compare $J(\mathbf{x})$ with $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$

$$J(\mathbf{x}) = f(\mathbf{x}) + \sum_{i=1}^{m} \mathbf{1}(g_i(\mathbf{x})), \text{ where } \mathbf{1}(z) = \left\{ \begin{array}{ll} 0 & \text{if } z \leq 0 \\ \infty & \text{otherwise} \end{array} \right.$$

v.s.

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \boldsymbol{\lambda}^\top \boldsymbol{g}(\mathbf{x}).$$

- $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is a lower bound of $J(\mathbf{x})$.
  $\therefore \ J(\mathbf{x}) = \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$

Compare $J(\mathbf{x})$ with $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$

$$J(\mathbf{x}) = f(\mathbf{x}) + \sum_{i=1}^{m} \mathbf{1}(g_i(\mathbf{x})), \text{ where } \mathbf{1}(z) = \begin{cases} 0 & \text{if } z \leq 0 \\ \infty & \text{otherwise} \end{cases}$$

v.s.

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \boldsymbol{\lambda}^{\top} \boldsymbol{g}(\mathbf{x}).$$

- $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is a lower bound of $J(\mathbf{x})$.
  $\therefore \ J(\mathbf{x}) = \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$.

- Recall the original problem:
  $$\min_{\mathbf{x} \in \mathbb{R}^d} \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$$

Compare $J(\mathbf{x})$ with $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$

$$J(\mathbf{x}) = f(\mathbf{x}) + \sum_{i=1}^{m} \mathbf{1}(g_i(\mathbf{x})), \text{ where } \mathbf{1}(z) = \left\{ \begin{array}{ll} 0 & \text{if } z \leq 0 \\ \infty & \text{otherwise} \end{array} \right.$$

v.s.

$$\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) + \boldsymbol{\lambda}^\top \mathbf{g}(\mathbf{x}).$$

- $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is a lower bound of $J(\mathbf{x})$.
  $\therefore$ $J(\mathbf{x}) = \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$.

- Recall the original problem:
  $$\min_{\mathbf{x} \in \mathbb{R}^d} \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$$

  Hence, $$\min_{\mathbf{x} \in \mathbb{R}^d} \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) \geq \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$$

Computation Issues

$$\min_{\mathbf{x} \in \mathbb{R}^d} \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) \geq \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$$

- $\min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is unconstrained, hence it is somehow easy to solve.

Computation Issues

$$\min_{\mathbf{x} \in \mathbb{R}^d} \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) \geq \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$$

- $\min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is unconstrained, hence it is somehow easy to solve.
- $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is affine w.r.t. $\boldsymbol{\lambda}$.

Computation Issues

$$\min_{\mathbf{x} \in \mathbb{R}^d} \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}) \geq \max_{\boldsymbol{\lambda} \geq \mathbf{0}} \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda}).$$

- $\min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is unconstrained, hence it is somehow easy to solve.
- $\mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$ is affine w.r.t. $\boldsymbol{\lambda}$.
- $\mathcal{D}(\boldsymbol{\lambda}) = \min_{\mathbf{x} \in \mathbb{R}^d} \mathcal{L}(\mathbf{x}, \boldsymbol{\lambda})$: pointwise minimum of affine functions of $\boldsymbol{\lambda}$.

Computation Issues

$$\min_{\mathbf{x}\in\mathbb{R}^d} \max_{\boldsymbol{\lambda}\geq\mathbf{0}} \mathcal{L}(\mathbf{x},\boldsymbol{\lambda}) \geq \max_{\boldsymbol{\lambda}\geq\mathbf{0}} \min_{\mathbf{x}\in\mathbb{R}^d} \mathcal{L}(\mathbf{x},\boldsymbol{\lambda}).$$

- $\min_{\mathbf{x}\in\mathbb{R}^d} \mathcal{L}(\mathbf{x},\boldsymbol{\lambda})$ is unconstrained, hence it is somehow easy to solve.

- $\mathcal{L}(\mathbf{x},\boldsymbol{\lambda})$ is affine w.r.t. $\boldsymbol{\lambda}$.

- $\mathcal{D}(\boldsymbol{\lambda}) = \min_{\mathbf{x}\in\mathbb{R}^d} \mathcal{L}(\mathbf{x},\boldsymbol{\lambda})$: pointwise minimum of affine functions of $\boldsymbol{\lambda}$.

  - The outer problem of maximization over $\boldsymbol{\lambda}$ can be efficiently computed.

    ($\mathcal{D}(\boldsymbol{\lambda})$ is concave so finding the maximum is easy).

Remark

- $h_j(\mathbf{x}) \geq 0 \iff -h_j(\mathbf{x}) \leq 0$.

Remark

- $h_j(\mathbf{x}) \geq 0 \iff -h_j(\mathbf{x}) \leq 0$.

- What's about "equality" constraints?

Remark

- $h_j(\mathbf{x}) \geq 0 \iff -h_j(\mathbf{x}) \leq 0$.

- What's about "equality" constraints?

$$h_j(\mathbf{x}) = 0 \implies \begin{cases} h_j(\mathbf{x}) \leq 0 & \text{and} \\ -h_j(\mathbf{x}) \leq 0. \end{cases}$$

Discussions