

Mathematics for Machine Learning

— Parameter Estimation

Joseph Chuang-Chieh Lin

Department of Computer Science & Engineering,
National Taiwan Ocean University

Fall 2025

Credits for the resource

- The slides are based on the textbooks:
 - *Marc Peter Deisenroth, A. Aldo Faisal, and Cheng Soon Ong: Mathematics for Machine Learning. Cambridge University Press. 2020.*
 - *Arnold J. Insel, Lawrence E. Spence, Stephen H. Friedberg: Linear Algebra, 4th Edition. Prentice Hall. 2013.*
 - *Howard Anton, Chris Rorres, Anton Kaul: Elementary Linear Algebra, 12th Edition. Wiley. 2019.*
- We could partially refer to the monograph:
Francesco Orabona: A Modern Introduction to Online Learning.
<https://arxiv.org/abs/1912.13213>

Outline

- 1 Maximum Likelihood Estimation
- 2 Maximum A Posteriori Estimation

Goal

- Use probabilistic distributions to model our uncertainty due to:
 - the observation process.
 - the uncertainty in the parameters of the predictor.

Outline

- 1 Maximum Likelihood Estimation
- 2 Maximum A Posteriori Estimation

Maximum Likelihood Estimation (MLE)

For data represented by a random variable \mathbf{x} and for a family of probability densities $p(\mathbf{x} \mid \boldsymbol{\theta})$ parameterized by $\boldsymbol{\theta}$, we aim at the **negative log-likelihood**:

$$\mathcal{L}_{\mathbf{x}}(\boldsymbol{\theta}) = -\log p(\mathbf{x} \mid \boldsymbol{\theta}).$$

- **Note:** The parameter $\boldsymbol{\theta}$ is varying and the data \mathbf{x} is fixed.
- $\mathcal{L}_{\mathbf{x}}(\boldsymbol{\theta})$: a function of $\boldsymbol{\theta}$.

Maximum Likelihood Estimation (MLE)

For data represented by a random variable \mathbf{x} and for a family of probability densities $p(\mathbf{x} \mid \boldsymbol{\theta})$ parameterized by $\boldsymbol{\theta}$, we aim at the **negative log-likelihood**:

$$\mathcal{L}_{\mathbf{x}}(\boldsymbol{\theta}) = -\log p(\mathbf{x} \mid \boldsymbol{\theta}).$$

- **Note:** The parameter $\boldsymbol{\theta}$ is varying and the data \mathbf{x} is fixed.
- $\mathcal{L}_{\mathbf{x}}(\boldsymbol{\theta})$: a function of $\boldsymbol{\theta}$.


For a given dataset \mathbf{x} , the likelihood allows us choose the settings of $\boldsymbol{\theta}$ that more “likely” has generated the data or how “likely” $\boldsymbol{\theta}$ is for the observations \mathbf{x} .

Example

- Specify that the conditional probability of the labels given the examples is a Gaussian distribution.
- Assume that we can explain our observation uncertainty by independent Gaussian noise $\varepsilon \sim \mathcal{N}(0, \sigma^2)$.
- We assume the linear model $\mathbf{x}_i^\top \boldsymbol{\theta}$ is used for prediction.

For each example-label pair (\mathbf{x}_i, y_i) ,

$$p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}) = \mathcal{N}(y_i \mid \mathbf{x}_i^\top \boldsymbol{\theta}, \sigma^2).$$



MLE for i.i.d. examples

- Assume that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ are i.i.d.
- The likelihood factorizes into a product of likelihoods of each individual example

$$p(\mathcal{Y} \mid \mathcal{X}, \boldsymbol{\theta}) = \prod_{i=1}^N p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}).$$

MLE for i.i.d. examples

- Assume that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ are i.i.d.
- The likelihood factorizes into a product of likelihoods of each individual example

$$p(\mathcal{Y} \mid \mathcal{X}, \boldsymbol{\theta}) = \prod_{i=1}^N p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}).$$

Then,

$$\mathcal{L}(\boldsymbol{\theta}) = -\log p(\mathcal{Y} \mid \mathcal{X}, \boldsymbol{\theta}) = -\sum_{i=1}^N \log p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}).$$

MLE for i.i.d. examples

- Assume that $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N)$ are i.i.d.
- The likelihood factorizes into a product of likelihoods of each individual example

$$p(\mathcal{Y} \mid \mathcal{X}, \boldsymbol{\theta}) = \prod_{i=1}^N p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}).$$

Then,

$$\mathcal{L}(\boldsymbol{\theta}) = -\log p(\mathcal{Y} \mid \mathcal{X}, \boldsymbol{\theta}) = -\sum_{i=1}^N \log p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}).$$

Note: Do not forget that $\mathcal{L}(\boldsymbol{\theta})$ is a function of $\boldsymbol{\theta}$.

Example (contd.)

$$\begin{aligned}\mathcal{L}(\boldsymbol{\theta}) &= -\sum_{i=1}^N \log p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}) = -\sum_{i=1}^N \log \mathcal{N}(y_i \mid \mathbf{x}_i^\top \boldsymbol{\theta}, \sigma^2) \\&= -\sum_{i=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \mathbf{x}_i^\top \boldsymbol{\theta})^2}{2\sigma^2}\right) \\&= -\sum_{i=1}^N \log \exp\left(-\frac{(y_i - \mathbf{x}_i^\top \boldsymbol{\theta})^2}{2\sigma^2}\right) - \sum_{i=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}} \\&= \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mathbf{x}_i^\top \boldsymbol{\theta})^2 - \sum_{i=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}}.\end{aligned}$$

\implies minimizing $\mathcal{L}(\boldsymbol{\theta})$

Example (contd.)

$$\begin{aligned}\mathcal{L}(\boldsymbol{\theta}) &= -\sum_{i=1}^N \log p(y_i \mid \mathbf{x}_i, \boldsymbol{\theta}) = -\sum_{i=1}^N \log \mathcal{N}(y_i \mid \mathbf{x}_i^\top \boldsymbol{\theta}, \sigma^2) \\&= -\sum_{i=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y_i - \mathbf{x}_i^\top \boldsymbol{\theta})^2}{2\sigma^2}\right) \\&= -\sum_{i=1}^N \log \exp\left(-\frac{(y_i - \mathbf{x}_i^\top \boldsymbol{\theta})^2}{2\sigma^2}\right) - \sum_{i=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}} \\&= \frac{1}{2\sigma^2} \sum_{i=1}^N (y_i - \mathbf{x}_i^\top \boldsymbol{\theta})^2 - \sum_{i=1}^N \log \frac{1}{\sqrt{2\pi\sigma^2}}.\end{aligned}$$

The second term is **constant**.

\implies minimizing $\mathcal{L}(\boldsymbol{\theta}) \implies$ solving the least-squares problem.

Issues

- This example has a closed-form solution.

Issues

- This example has a **closed-form solution**.
 - No closed-form \Rightarrow resort to numerical optimization.
- MLE may suffer from overfitting

Issues

- This example has a **closed-form solution**.
 - No closed-form \Rightarrow resort to numerical optimization.
- MLE may suffer from overfitting (analogous to unregularized empirical risk minimization).

Outline

- 1 Maximum Likelihood Estimation
- 2 Maximum A Posteriori Estimation

Motivation (1/2)

What if we have **prior knowledge** about the distribution of the **parameters θ** ?

Motivation (1/2)

What if we have **prior knowledge** about the distribution of the **parameters θ** ?

We can **multiply an additional term (i.e., $p(\theta)$)** to the likelihood.

Motivation (2/2)

- For a given prior, **after observing some data \mathbf{x}** , how should we update $p(\theta)$?
 - \Rightarrow Bayes's theorem.
 - ★ Compute a posterior distribution $p(\theta | \mathbf{x})$.

$$p(\theta | \mathbf{x}) = \frac{p(\mathbf{x} | \theta)p(\theta)}{p(\mathbf{x})}.$$

Motivation (2/2)


- For a given prior, **after observing some data \mathbf{x}** , how should we update $p(\theta)$?
 - ⇒ Bayes's theorem.
 - ★ Compute a posterior distribution $p(\theta | \mathbf{x})$.


$$p(\theta | \mathbf{x}) = \frac{p(\mathbf{x} | \theta)p(\theta)}{p(\mathbf{x})}.$$

So,


$$p(\theta | \mathbf{x}) \propto p(\mathbf{x} | \theta)p(\theta).$$

MLE vs. MAP






(a) Overfitting



(b) Underfitting.



(c) Fitting well.

Discussions