

Mathematics for Machine Learning

— Vector Calculus: Gradients of Vector-Valued Functions and Matrices

Joseph Chuang-Chieh Lin

Department of Computer Science & Information Engineering,
Tamkang University

Fall 2023

Credits for the resource

- The slides are based on the textbooks:
 - *Marc Peter Deisenroth, A. Aldo Faisal, and Cheng Soon Ong: Mathematics for Machine Learning. Cambridge University Press. 2020.*
 - *Howard Anton, Chris Rorres, Anton Kaul: Elementary Linear Algebra. Wiley. 2019.*
- We could partially refer to the monograph:
Francesco Orabona: A Modern Introduction to Online Learning.
<https://arxiv.org/abs/1912.13213>

Outline

1 Gradients of Vector-Valued Functions

2 Gradients of Matrices

Outline

1 Gradients of Vector-Valued Functions

2 Gradients of Matrices

Our Focus

- Partial derivatives and gradients of functions $f : \mathbb{R}^n \mapsto \mathbb{R}^m$, for $n \geq 1, m > 1$.

Vector of Functions

Given

- $\mathbf{f} : \mathbb{R}^n \mapsto \mathbb{R}^m$.
- $\mathbf{x} = [x_1, \dots, x_n]^\top \in \mathbb{R}^n$.

The corresponding *vector of functions*:

$$\mathbf{f}[\mathbf{x}] = \begin{bmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{bmatrix} \in \mathbb{R}^m.$$

Vector of Functions

Given

- $\mathbf{f} : \mathbb{R}^n \mapsto \mathbb{R}^m$.
- $\mathbf{x} = [x_1, \dots, x_n]^\top \in \mathbb{R}^n$.

The corresponding *vector of functions*:

$$\mathbf{f}[\mathbf{x}] = \begin{bmatrix} f_1(\mathbf{x}) \\ \vdots \\ f_m(\mathbf{x}) \end{bmatrix} \in \mathbb{R}^m.$$

We can view \mathbf{f} as $[f_1, \dots, f_m]^\top$, such that $f_i : \mathbb{R}^n \mapsto \mathbb{R}$.

Therefore,

$$\frac{\partial \mathbf{f}}{\partial x_i} = \begin{bmatrix} \frac{\partial f_1}{\partial x_i} \\ \vdots \\ \frac{\partial f_m}{\partial x_i} \end{bmatrix} = \begin{bmatrix} \lim_{h \rightarrow 0} \frac{f_1(x_1, \dots, x_{i-1}, x_i + h, x_{i+1}, \dots, x_n) - f_1(\mathbf{x})}{h} \\ \vdots \\ \lim_{h \rightarrow 0} \frac{f_m(x_1, \dots, x_{i-1}, x_i + h, x_{i+1}, \dots, x_n) - f_m(\mathbf{x})}{h} \end{bmatrix} \in \mathbb{R}^m.$$

So,

$$\begin{aligned}\frac{d\mathbf{f}}{d\mathbf{x}} &= \left[\frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_1} \cdots \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_n} \right] \\ &= \begin{bmatrix} \frac{\partial f_1(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_1(\mathbf{x})}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_m(\mathbf{x})}{\partial x_n} \end{bmatrix} \in \mathbb{R}^{m \times n}.\end{aligned}$$

So,

$$\begin{aligned}\frac{d\mathbf{f}}{d\mathbf{x}} &= \left[\frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_1} \cdots \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_n} \right] \\ &= \begin{bmatrix} \frac{\partial f_1(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_1(\mathbf{x})}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_m(\mathbf{x})}{\partial x_n} \end{bmatrix} \in \mathbb{R}^{m \times n}.\end{aligned}$$

- We call this collection of all first-order partial derivatives of a vector-valued function \mathbf{f} the **Jacobian**.

So,

$$\begin{aligned}\frac{d\mathbf{f}}{d\mathbf{x}} &= \left[\frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_1} \cdots \frac{\partial \mathbf{f}(\mathbf{x})}{\partial x_n} \right] \\ &= \begin{bmatrix} \frac{\partial f_1(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_1(\mathbf{x})}{\partial x_n} \\ \vdots & & \vdots \\ \frac{\partial f_m(\mathbf{x})}{\partial x_1} & \cdots & \frac{\partial f_m(\mathbf{x})}{\partial x_n} \end{bmatrix} \in \mathbb{R}^{m \times n}.\end{aligned}$$

- We call this collection of all first-order partial derivatives of a vector-valued function \mathbf{f} the **Jacobian**.

- ★ Denote by $\mathbf{J} = \nabla_{\mathbf{x}} \mathbf{f} = \frac{d\mathbf{f}(\mathbf{x})}{d\mathbf{x}}$
 - $J(i, j) = \frac{\partial f_i}{\partial x_j}$.

Example

Derivative of a Polynomial

Given $\mathbf{f}(\mathbf{x}) = \mathbf{A}\mathbf{x}$, $\mathbf{f}(\mathbf{x}) \in \mathbb{R}^M$, $\mathbf{A} \in \mathbb{R}^{M \times N}$, and $\mathbf{x} \in \mathbb{R}^N$. Compute

$$\frac{d\mathbf{f}}{d\mathbf{x}} = ?$$

Note:

- $\mathbf{f} : \mathbb{R}^N \mapsto \mathbb{R}^M$, so $\frac{d\mathbf{f}}{d\mathbf{x}} \in \mathbb{R}^{M \times N}$.
- $f_i(\mathbf{x}) =$

Example

Derivative of a Polynomial

Given $\mathbf{f}(\mathbf{x}) = \mathbf{A}\mathbf{x}$, $\mathbf{f}(\mathbf{x}) \in \mathbb{R}^M$, $\mathbf{A} \in \mathbb{R}^{M \times N}$, and $\mathbf{x} \in \mathbb{R}^N$. Compute

$$\frac{d\mathbf{f}}{d\mathbf{x}} = ?$$

Note:

- $\mathbf{f} : \mathbb{R}^N \mapsto \mathbb{R}^M$, so $\frac{d\mathbf{f}}{d\mathbf{x}} \in \mathbb{R}^{M \times N}$.
- $f_i(\mathbf{x}) = \sum_{j=1}^N A_{ij}x_j$

Example

Derivative of a Polynomial

Given $\mathbf{f}(\mathbf{x}) = \mathbf{A}\mathbf{x}$, $\mathbf{f}(\mathbf{x}) \in \mathbb{R}^M$, $\mathbf{A} \in \mathbb{R}^{M \times N}$, and $\mathbf{x} \in \mathbb{R}^N$. Compute

$$\frac{d\mathbf{f}}{d\mathbf{x}} = ?$$

Note:

- $\mathbf{f} : \mathbb{R}^N \mapsto \mathbb{R}^M$, so $\frac{d\mathbf{f}}{d\mathbf{x}} \in \mathbb{R}^{M \times N}$.
- $f_i(\mathbf{x}) = \sum_{j=1}^N A_{ij}x_j \implies \frac{\partial f_i}{\partial x_j} = A_{ij}$.

Example

Derivative of a Polynomial

Given $\mathbf{f}(\mathbf{x}) = \mathbf{A}\mathbf{x}$, $\mathbf{f}(\mathbf{x}) \in \mathbb{R}^M$, $\mathbf{A} \in \mathbb{R}^{M \times N}$, and $\mathbf{x} \in \mathbb{R}^N$. Compute

$$\frac{d\mathbf{f}}{d\mathbf{x}} = \mathbf{A}$$

Note:

- $\mathbf{f} : \mathbb{R}^N \mapsto \mathbb{R}^M$, so $\frac{d\mathbf{f}}{d\mathbf{x}} \in \mathbb{R}^{M \times N}$.
- $f_i(\mathbf{x}) = \sum_{j=1}^N A_{ij}x_j \implies \frac{\partial f_i}{\partial x_j} = A_{ij}$.

Example: Gradient of a Least-Squared Loss in a Linear Model

Consider the linear model

$$\mathbf{y} = \Phi\boldsymbol{\theta},$$

where

- $\boldsymbol{\theta} \in \mathbb{R}^D$: a parameter vector
- $\Phi \in \mathbb{R}^{N \times D}$: input features
- $\mathbf{y} \in \mathbb{R}^N$: the corresponding observations.

We define that

$$L(\mathbf{e}) := \|\mathbf{e}\|^2.$$

$$\mathbf{e}(\boldsymbol{\theta}) := \mathbf{y} - \Phi\boldsymbol{\theta}.$$

Compute $\frac{\partial L}{\partial \boldsymbol{\theta}}$ (using the chain rule).

Example (2/3)

Note that

- $\frac{\partial L}{\partial \theta} \in \mathbb{R}^{1 \times D}$

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D} \quad (\because L : \mathbb{R}^D \mapsto \mathbb{R}).$
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D}$ ($\because L : \mathbb{R}^D \mapsto \mathbb{R}$).
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$ ($\because \mathbf{e} : \mathbb{R}^D \mapsto \mathbb{R}^N$).

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D}$ ($\because L : \mathbb{R}^D \mapsto \mathbb{R}$).
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$ ($\because \mathbf{e} : \mathbb{R}^D \mapsto \mathbb{R}^N$).
- $\frac{\partial L}{\partial \mathbf{e}} \in \mathbb{R}^{1 \times N}$

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D}$ ($\because L : \mathbb{R}^D \mapsto \mathbb{R}$).
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$ ($\because \mathbf{e} : \mathbb{R}^D \mapsto \mathbb{R}^N$).
- $\frac{\partial L}{\partial \mathbf{e}} \in \mathbb{R}^{1 \times N}$ ($\because L : \mathbb{R}^N \mapsto \mathbb{R}$).

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D}$ ($\because L : \mathbb{R}^D \mapsto \mathbb{R}$).
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$ ($\because \mathbf{e} : \mathbb{R}^D \mapsto \mathbb{R}^N$).
- $\frac{\partial L}{\partial \mathbf{e}} \in \mathbb{R}^{1 \times N}$ ($\because L : \mathbb{R}^N \mapsto \mathbb{R}$).
- $\frac{\partial L}{\partial \boldsymbol{\theta}} = \frac{\partial L}{\partial \mathbf{e}} \frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}}$ (chain rule).

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D}$ ($\because L : \mathbb{R}^D \mapsto \mathbb{R}$).
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$ ($\because \mathbf{e} : \mathbb{R}^D \mapsto \mathbb{R}^N$).
- $\frac{\partial L}{\partial \mathbf{e}} \in \mathbb{R}^{1 \times N}$ ($\because L : \mathbb{R}^N \mapsto \mathbb{R}$).
- $\frac{\partial L}{\partial \boldsymbol{\theta}} = \frac{\partial L}{\partial \mathbf{e}} \frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}}$ (chain rule).
- The d th element: $\frac{\partial L}{\partial \boldsymbol{\theta}}[1, d] = \sum_{i=1}^N \frac{\partial L}{\partial \mathbf{e}}[i] \frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}}[i, d]$.

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D}$ ($\because L : \mathbb{R}^D \mapsto \mathbb{R}$).
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$ ($\because \mathbf{e} : \mathbb{R}^D \mapsto \mathbb{R}^N$).
- $\frac{\partial L}{\partial \mathbf{e}} \in \mathbb{R}^{1 \times N}$ ($\because L : \mathbb{R}^N \mapsto \mathbb{R}$).
- $\frac{\partial L}{\partial \boldsymbol{\theta}} = \frac{\partial L}{\partial \mathbf{e}} \frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}}$ (chain rule).
- The d th element: $\frac{\partial L}{\partial \boldsymbol{\theta}}[1, d] = \sum_{i=1}^N \frac{\partial L}{\partial \mathbf{e}}[i] \frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}}[i, d]$.
- $L = \|\mathbf{e}\|^2 = \mathbf{e}^\top \mathbf{e}$ and $\frac{\partial L}{\partial \mathbf{e}} = 2\mathbf{e}^\top \in \mathbb{R}^{1 \times N}$.

Example (2/3)

Note that

- $\frac{\partial L}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{1 \times D}$ ($\because L : \mathbb{R}^D \mapsto \mathbb{R}$).
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} \in \mathbb{R}^{N \times D}$ ($\because \mathbf{e} : \mathbb{R}^D \mapsto \mathbb{R}^N$).
- $\frac{\partial L}{\partial \mathbf{e}} \in \mathbb{R}^{1 \times N}$ ($\because L : \mathbb{R}^N \mapsto \mathbb{R}$).
- $\frac{\partial L}{\partial \boldsymbol{\theta}} = \frac{\partial L}{\partial \mathbf{e}} \frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}}$ (chain rule).
- The d th element: $\frac{\partial L}{\partial \boldsymbol{\theta}}[1, d] = \sum_{i=1}^N \frac{\partial L}{\partial \mathbf{e}}[i] \frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}}[i, d]$.
- $L = \|\mathbf{e}\|^2 = \mathbf{e}^\top \mathbf{e}$ and $\frac{\partial L}{\partial \mathbf{e}} = 2\mathbf{e}^\top \in \mathbb{R}^{1 \times N}$.
- $\frac{\partial \mathbf{e}}{\partial \boldsymbol{\theta}} = -\Phi \in \mathbb{R}^{N \times D}$.

Example (3/3)

$$\frac{\partial L}{\partial \boldsymbol{\theta}} = -2\mathbf{e}^\top \boldsymbol{\Phi} = -2(\mathbf{y}^\top - \boldsymbol{\theta}^\top \boldsymbol{\Phi}^\top) \boldsymbol{\Phi} \in \mathbb{R}^{1 \times D}$$

Example (3/3)

$$\frac{\partial L}{\partial \boldsymbol{\theta}} = -2\mathbf{e}^\top \Phi = -2(\mathbf{y}^\top - \boldsymbol{\theta}^\top \Phi^\top) \Phi \in \mathbb{R}^{1 \times D}$$

By the way, we can obtain the same result without using the chain rule:

$$L_2(\boldsymbol{\theta}) := \|\mathbf{y} - \Phi \boldsymbol{\theta}\|^2 = (\mathbf{y} - \Phi \boldsymbol{\theta})^\top (\mathbf{y} - \Phi \boldsymbol{\theta}).$$

Example (3/3)

$$\frac{\partial L}{\partial \boldsymbol{\theta}} = -2\mathbf{e}^\top \Phi = -2(\mathbf{y}^\top - \boldsymbol{\theta}^\top \Phi^\top) \Phi \in \mathbb{R}^{1 \times D}$$

By the way, we can obtain the same result without using the chain rule:

$$L_2(\boldsymbol{\theta}) := \|\mathbf{y} - \Phi \boldsymbol{\theta}\|^2 = (\mathbf{y} - \Phi \boldsymbol{\theta})^\top (\mathbf{y} - \Phi \boldsymbol{\theta}).$$

- It becomes impractical for deep function compositions.

Outline

1 Gradients of Vector-Valued Functions

2 Gradients of Matrices

Motivations

- There are scenarios that we need to take gradients of matrices w.r.t. vectors (or other matrices).
 - ⇒ This results in a multidimensional tensor.
 - Multidimensional array.
- Compute the gradient of an $m \times n$ matrix \mathbf{A} w.r.t. a $p \times q$ matrix \mathbf{B} :
 - The Jacobian \mathbf{J} would be $(m \times n) \times (p \times q)$ (4-dimensional tensor).
 - $J_{ijkl} = \frac{\partial A_{ij}}{\partial B_{kl}}$.
- Matrices \Leftrightarrow linear mappings, so


Motivations

- There are scenarios that we need to take gradients of matrices w.r.t. vectors (or other matrices).
 - ⇒ This results in a multidimensional tensor.
 - Multidimensional array.
- Compute the gradient of an $m \times n$ matrix \mathbf{A} w.r.t. a $p \times q$ matrix \mathbf{B} :
 - The Jacobian \mathbf{J} would be $(m \times n) \times (p \times q)$ (4-dimensional tensor).
 - $J_{ijkl} = \frac{\partial A_{ij}}{\partial B_{kl}}$.
- Matrices \Leftrightarrow linear mappings, so

There is a vector-space isomorphism (i.e., linear, invertible mapping) between the space $\mathbb{R}^{m \times n}$ of $m \times n$ matrices and the space \mathbb{R}^{mn} of mn vector.

Visualization of Two Approaches for the Isomorphism


$$\mathbf{A} \in \mathbb{R}^{4 \times 2}$$



$$\mathbf{x} \in \mathbb{R}^3$$



Partial derivatives:



Visualization of Two Approaches for the Isomorphism


$$\mathbf{A} \in \mathbb{R}^{4 \times 2}$$



$$\mathbf{x} \in \mathbb{R}^3$$



$$\mathbf{A} \in \mathbb{R}^{4 \times 2}$$




re-shape

$$\tilde{\mathbf{A}} \in \mathbb{R}^8$$




gradient

$$\frac{d\tilde{\mathbf{A}}}{d\mathbf{x}} \in \mathbb{R}^{8 \times 3}$$



re-shape

$$\frac{d\mathbf{A}}{d\mathbf{x}} \in \mathbb{R}^{4 \times 2 \times 3}$$



Example: Gradient of Vectors w.r.t. Matrices

Consider

$$\mathbf{f} = \mathbf{A}\mathbf{x}, \text{ where } \mathbf{f} \in \mathbb{R}^M, \mathbf{A} \in \mathbb{R}^{M \times N}, \mathbf{x} \in \mathbb{R}^N.$$

Goal: Compute the gradient $\frac{d\mathbf{f}}{d\mathbf{A}}$.

Note:

- $\frac{d\mathbf{f}}{d\mathbf{A}} \in$

Example: Gradient of Vectors w.r.t. Matrices

Consider

$$\mathbf{f} = \mathbf{A}\mathbf{x}, \text{ where } \mathbf{f} \in \mathbb{R}^M, \mathbf{A} \in \mathbb{R}^{M \times N}, \mathbf{x} \in \mathbb{R}^N.$$

Goal: Compute the gradient $\frac{d\mathbf{f}}{d\mathbf{A}}$.

Note:

- $\frac{d\mathbf{f}}{d\mathbf{A}} \in \mathbb{R}^{M \times (M \times N)}$

$$\frac{d\mathbf{f}}{d\mathbf{A}} =$$

$$\frac{d\mathbf{f}}{d\mathbf{A}} = \begin{bmatrix} \frac{\partial f_1}{\partial \mathbf{A}} \\ \vdots \\ \frac{\partial f_M}{\partial \mathbf{A}} \end{bmatrix},$$

$$\frac{d\mathbf{f}}{d\mathbf{A}} = \begin{bmatrix} \frac{\partial f_1}{\partial \mathbf{A}} \\ \vdots \\ \frac{\partial f_M}{\partial \mathbf{A}} \end{bmatrix}, \frac{\partial f_i}{\partial \mathbf{A}} \in \mathbb{R}^{1 \times (M \times N)}.$$

$$\frac{d\mathbf{f}}{d\mathbf{A}} = \begin{bmatrix} \frac{\partial f_1}{\partial \mathbf{A}} \\ \vdots \\ \frac{\partial f_M}{\partial \mathbf{A}} \end{bmatrix}, \quad \frac{\partial f_i}{\partial \mathbf{A}} \in \mathbb{R}^{1 \times (M \times N)}.$$

- We can explicitly expand $f_i = \sum_{j=1}^N A_{ij}x_j$, for $i = 1, \dots, M$.

Hence,

$$\frac{\partial f_i}{\partial A_{iq}} = x_q.$$

So we can derive

$$\frac{\partial f_i}{\partial A_{i,:}} = \mathbf{x}^\top$$

$$\frac{d\mathbf{f}}{d\mathbf{A}} = \begin{bmatrix} \frac{\partial f_1}{\partial \mathbf{A}} \\ \vdots \\ \frac{\partial f_M}{\partial \mathbf{A}} \end{bmatrix}, \quad \frac{\partial f_i}{\partial \mathbf{A}} \in \mathbb{R}^{1 \times (M \times N)}.$$

- We can explicitly expand $f_i = \sum_{j=1}^N A_{ij}x_j$, for $i = 1, \dots, M$.

Hence,

$$\frac{\partial f_i}{\partial A_{iq}} = x_q.$$

So we can derive

$$\frac{\partial f_i}{\partial A_{i,:}} = \mathbf{x}^\top \in \mathbb{R}^{1 \times (1 \times N)} \quad \text{and} \quad \frac{\partial f_i}{\partial A_{k \neq i,:}} = \mathbf{0}^\top \in \mathbb{R}^{1 \times (1 \times N)}.$$

Stack the partial derivatives:

$$\frac{\partial f_i}{\partial \mathbf{A}} = \begin{bmatrix} \mathbf{0}^\top \\ \vdots \\ \mathbf{0}^\top \\ \mathbf{x}^\top \\ \mathbf{0}^\top \\ \vdots \\ \mathbf{0}^\top \end{bmatrix} \in \mathbb{R}^{1 \times (M \times N)}$$

Example: Gradient of Matrices w.r.t. Matrices

Consider a matrix $\mathbf{R} \in \mathbb{R}^{M \times N}$ and $\mathbf{f} : \mathbb{R}^{M \times N} \mapsto \mathbb{R}^{N \times N}$ with

$$\mathbf{f}(\mathbf{R}) = \mathbf{R}^\top \mathbf{R} := \mathbf{K} \in \mathbb{R}^{N \times N}$$

Goal: Compute the gradient $\frac{d\mathbf{K}}{d\mathbf{R}}$.

Note:

- $\frac{d\mathbf{K}}{d\mathbf{R}} \in$

Example: Gradient of Matrices w.r.t. Matrices

Consider a matrix $\mathbf{R} \in \mathbb{R}^{M \times N}$ and $\mathbf{f} : \mathbb{R}^{M \times N} \mapsto \mathbb{R}^{N \times N}$ with

$$\mathbf{f}(\mathbf{R}) = \mathbf{R}^\top \mathbf{R} := \mathbf{K} \in \mathbb{R}^{N \times N}$$

Goal: Compute the gradient $\frac{d\mathbf{K}}{d\mathbf{R}}$.

Note:

- $\frac{d\mathbf{K}}{d\mathbf{R}} \in \mathbb{R}^{(N \times N) \times (M \times N)}$.
- $\frac{dK_{pq}}{d\mathbf{R}} \in \mathbb{R}^{1 \times (M \times N)}$, for $p, q = 1, \dots, N$, K_{pq} : the (p, q) th entry of \mathbf{K} .

$$K_{pq} = \mathbf{r}_p^\top \mathbf{r}_q = \sum_{t=1}^M R_{tp} R_{tq}.$$

\mathbf{r}_i : the i th column of \mathbf{R} .

Example (2/2)

Compute $\frac{\partial K_{pq}}{\partial R_{ij}}$: (sum rule)

$$\frac{\partial K_{pq}}{\partial R_{ij}} = \sum_{t=1}^M \frac{\partial}{\partial R_{ij}} R_{tp} R_{tq} = \partial_{pqij},$$

where

$$\partial_{pqij} = \begin{cases} R_{iq} & \text{if } j = p, p \neq q \\ R_{ip} & \text{if } j = q, p \neq q \\ 2R_{iq} & \text{if } j = p, p = q \\ 0 & \text{otherwise} \end{cases}$$

Example (2/2)

Compute $\frac{\partial K_{pq}}{\partial R_{ij}}$: (sum rule)

$$\frac{\partial K_{pq}}{\partial R_{ij}} = \sum_{t=1}^M \frac{\partial}{\partial R_{ij}} R_{tp} R_{tq} = \partial_{pqij},$$

where

$$\partial_{pqij} = \begin{cases} R_{iq} & \text{if } j = p, p \neq q \\ R_{ip} & \text{if } j = q, p \neq q \\ 2R_{iq} & \text{if } j = p, p = q \\ 0 & \text{otherwise} \end{cases}$$

Hence, each entry of the desired gradient $\frac{d\mathbf{K}}{d\mathbf{R}} \in \mathbb{R}^{(N \times N) \times (M \times N)}$ is ∂_{pqij} , for $p, q, j = 1, \dots, N$ and $i = 1, \dots, M$.

Useful Identities for Computing Gradients (1/2)

$$\frac{\partial}{\partial \mathbf{X}} \mathbf{f}(\mathbf{X})^\top = \left(\frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right)^\top.$$

$$\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{f}(\mathbf{X})) = \text{tr} \left(\frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right).$$

$$\frac{\partial}{\partial \mathbf{X}} \det(\mathbf{f}(\mathbf{X})) = \det(\mathbf{f}(\mathbf{X})) \text{tr} \left(\mathbf{f}(\mathbf{X})^{-1} \frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right)$$

Useful Identities for Computing Gradients (1/2)

$$\frac{\partial}{\partial \mathbf{X}} \mathbf{f}(\mathbf{X})^\top = \left(\frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right)^\top.$$

$$\frac{\partial}{\partial \mathbf{X}} \text{tr}(\mathbf{f}(\mathbf{X})) = \text{tr} \left(\frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right).$$

$$\frac{\partial}{\partial \mathbf{X}} \det(\mathbf{f}(\mathbf{X})) = \det(\mathbf{f}(\mathbf{X})) \text{tr} \left(\mathbf{f}(\mathbf{X})^{-1} \frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \right) \implies \text{Exercise}$$

$$\frac{\partial}{\partial \mathbf{X}} \mathbf{f}(\mathbf{X})^{-1} = -\mathbf{f}(\mathbf{X})^{-1} \frac{\partial \mathbf{f}(\mathbf{X})}{\partial \mathbf{X}} \mathbf{f}(\mathbf{X})^{-1}$$

$$\frac{\partial \mathbf{a}^\top \mathbf{X}^{-1} \mathbf{b}}{\partial \mathbf{X}} = -(\mathbf{X}^{-1})^\top \mathbf{a} \mathbf{b}^\top (\mathbf{X}^{-1})^\top$$

Useful Identities for Computing Gradients (2/2)

$$\frac{\partial \mathbf{x}^\top \mathbf{a}}{\partial \mathbf{x}} = \mathbf{a}^\top$$

$$\frac{\partial \mathbf{a}^\top \mathbf{x}}{\partial \mathbf{x}} = \mathbf{a}^\top$$

$$\frac{\partial \mathbf{a}^\top \mathbf{X} \mathbf{b}}{\partial \mathbf{X}} = \mathbf{a} \mathbf{b}^\top$$

$$\frac{\partial \mathbf{x}^\top \mathbf{B} \mathbf{x}}{\partial \mathbf{x}} = \mathbf{x}^\top (\mathbf{B} + \mathbf{B}^\top)$$

$$\frac{\partial}{\partial \mathbf{s}} (\mathbf{x} - \mathbf{A} \mathbf{s})^\top \mathbf{W} (\mathbf{x} - \mathbf{A} \mathbf{s}) = -2(\mathbf{x} - \mathbf{A} \mathbf{s})^\top \mathbf{W} \mathbf{A} \text{ for symmetric } \mathbf{W}.$$

Discussions