

Randomized Algorithms

The Lovász Local Lemma

Joseph Chuang-Chieh Lin

Department of Computer Science & Engineering,
National Taiwan Ocean University

Spring 2026



Outline

- 1 The Lovász Local Lemma
- 2 Applications
- 3 Explicit Constructions Using the Local Lemma
- 4 Lovász Local Lemma: The General Case



Outline

- 1 The Lovász Local Lemma
- 2 Applications
- 3 Explicit Constructions Using the Local Lemma
- 4 Lovász Local Lemma: The General Case



Introduction

- One of the most elegant and useful tools in applying the probabilistic method is the *Lovász Local Lemma* (LLL).
- Let $E_1, E_2, \dots, E_n$ be a set of **bad** events, we want to show that there is an event (or element in the sample space) that is NOT included in any of the bad events.



Mutually Independent (Recall)

Events $E_1, E_2, \dots, E_n$ are mutually independent iff for any subset $I \subseteq [1, n]$,

$$\Pr\left[\bigcap_{i \in I} E_i\right] = \prod_{i \in I} \Pr[E_i].$$

- If $E_1, E_2, \dots, E_n$ are mutually independent then so are $\bar{E}_1, \bar{E}_2, \dots, \bar{E}_n$. (Exercise)
- If $\Pr[E_i] < 1$ for all i , then

$$\Pr\left[\bigcap_{i=1}^n \bar{E}_i\right] = \prod_{i=1}^n \Pr[\bar{E}_i] > 0.$$

- In many probabilistic-method proofs, requiring all bad events to be mutually independent is an **unrealistically strong assumption**.



LLL Comes to the Rescue

- The Lovász local lemma generalizes to the case where the n events are **not mutually independent** but **the dependency is limited**.
- Specifically, we say that an event E_{n+1} is mutually independent of the events $E_1, E_2, \dots, E_n$ if, for any subset $I \subseteq [1, n]$,

$$\Pr[E_{n+1} \mid \bigcap_{j \in I} E_j] = \Pr[E_{n+1}]$$



Dependency graph

Dependency graph

A dependency graph for a set of events $E_1, E_2, \dots, E_n$ is a graph $G = (V, E)$ such that $V = \{1, 2, \dots, n\}$ and for $i = 1, 2, \dots, n$, event E_i is mutually independent of the events $\{E_j \mid (i, j) \notin E\}$.

- We discuss first a special case, the symmetric version of the Lovász local lemma, which is more intuitive and sufficient for most algorithmic applications.


Example of a Dependency Graph

Let X_1, X_2, X_3, X_4 be four independent fair coin flips.

$$E_1 = \{X_1 = H \text{ and } X_2 = H\},$$

$$E_2 = \{X_2 = H \text{ and } X_3 = H\},$$

$$E_3 = \{X_3 = H \text{ and } X_4 = H\}.$$



- E_1 and E_2 are connected because both depend on X_2 .
- E_2 and E_3 are connected because both depend on X_3 .
- E_1 and E_3 are not connected because they depend on disjoint coin flips.



General Idea of the Lovász Local Lemma

If a set of “bad” events

- are *mostly* mutually independent, and
- happen with *low* probability

then with positive probability none of them happen.

Lovász Local Lemma

Theorem [Lovász Local Lemma, Lovász & Erdős 1975]

Let $E_1, E_2, \dots, E_n$ be a set of events, and assume that the following hold:

- 1 for all i , $\Pr[E_i] \leq p$;
- 2 the degree of the dependency graph given by $E_1, E_2, \dots, E_n$ is bounded by d ;
- 3 $4dp \leq 1$.

Then

$$\Pr\left[\bigcap_{i=1}^n \bar{E}_i\right] > 0.$$

- Note that $p < 1$ for $d > 0$ ($d = 0$: the mutually independent case).



Proof of LLL (1/9)

- Let $S \subset \{1, 2, \dots, n\}$. We prove by induction on $s = 0, 1, \dots, n-1$ that, if $|S| \leq s$, then for all $k \notin S$ we have

$$\Pr[E_k \mid \bigcap_{j \in S} \bar{E}_j] \leq 2p.$$

- For this expression to be well-defined when $S \neq \emptyset$, we need

$$\Pr\left[\bigcap_{j \in S} \bar{E}_j\right] > 0.$$

*Note that $\Pr[A \mid B] = \Pr[A \cap B] / \Pr[B]$.



Proof of LLL (2/9)

- The base case $s = 0$ follows from the assumption that $\Pr[E_k] \leq p$.
- The inductive step:
 - First, show that $\Pr \left[\bigcap_{j \in S} \bar{E}_j \right] > 0$.

This is true when $s = 1$, because $\Pr[\bar{E}_j] \geq 1 - p > 0$.

For $s > 1$, WLOG, $S = \{1, 2, \dots, s\}$. Then ($\because 4dp \leq 1$)

$$\begin{aligned}
 \Pr \left[\bigcap_{i=1}^s \bar{E}_i \right] &= \prod_{i=1}^s \Pr \left[\bar{E}_i \mid \bigcap_{j=1}^{i-1} \bar{E}_j \right] \\
 &= \prod_{i=1}^s \left(1 - \Pr \left[E_i \mid \bigcap_{j=1}^{i-1} \bar{E}_j \right] \right) \\
 &\geq \prod_{i=1}^s (1 - 2p) > 0.
 \end{aligned}$$



Proof of LLL (3/9)

- For the rest of the induction, let $S_1 = \{j \in S \mid (k, j) \in E\}$ and $S_2 = S \setminus S_1$. (note: $k \notin S$ by assumption)
 - If $S_2 = S$ then E_k is mutually independent of the events $\bar{E}_i, i \in S$, and

$$\Pr[E_k \mid \bigcap_{j \in S} \bar{E}_j] = \Pr[E_k] \leq p.$$

- We continue with the case $|S_2| < s$.



Proof of LLL (4/9)

- It will be helpful to introduce the following notation.
- Let F_S be defined by

$$F_S = \bigcap_{j \in S} \overline{E_j},$$

And similarly define F_{S_1} and F_{S_2} . Notice that $F_S = F_{S_1} \cap F_{S_2}$.

- By the definition of conditional probability,

$$\Pr[E_k \mid F_S] = \frac{\Pr[E_k \cap F_S]}{\Pr[F_S]}. \quad (*)$$



Proof of LLL (6/9)

$$\Pr[E_k \mid F_S] = \frac{\Pr[E_k \cap F_S]}{\Pr[F_S]}. \quad (*)$$

- Applying the definition of conditional probability to (*), we obtain

$$\Pr[E_k \cap F_S] = \Pr[E_k \cap F_{S_1} \cap F_{S_2}] = \Pr[E_k \cap F_{S_1} \mid F_{S_2}] \Pr[F_{S_2}].$$

- The denominator can be written as

$$\Pr[F_S] = \Pr[F_{S_1} \cap F_{S_2}] = \Pr[F_{S_1} \mid F_{S_2}] \Pr[F_{S_2}].$$

- Canceling the common factor (nonzero), yields

$$\Pr[E_k \mid F_S] = \frac{\Pr[E_k \cap F_{S_1} \mid F_{S_2}]}{\Pr[F_{S_1} \mid F_{S_2}]}. \quad (**)$$



Proof of LLL (7/9)

- Since
 - the probability of an intersection of events is bounded by the probability of any one of the events, and
 - E_k is independent of the events in S_2 ,

we can bound the numerator of (**) by

$$\Pr[E_k \cap F_{S_1} \mid F_{S_2}] \leq \Pr[E_k \mid F_{S_2}] = \Pr[E_k] \leq p.$$

Because $|S_2| < |S| = s$, we can apply the induction hypothesis to

$$\Pr[E_i \mid F_{S_2}] = \Pr\left[E_i \mid \bigcap_{j \in S_2} \overline{E_j}\right].$$



Proof of LLL (8/9)

- Using also the fact that $|S_1| \leq d$, we establish a lower bound on the denominator of $(**)$ as follows:

$$\begin{aligned}\Pr[F_{S_1} \mid F_{S_2}] &= \Pr \left[\bigcap_{i \in S_1} \bar{E}_i \mid \bigcap_{j \in S_2} \bar{E}_j \right] \\ &\geq 1 - \sum_{i \in S_1} \Pr \left[E_i \mid \bigcap_{j \in S_2} \bar{E}_j \right] \\ &\geq 1 - \sum_{i \in S_1} 2p \\ &\geq 1 - 2pd \\ &\geq \frac{1}{2}.\end{aligned}$$



Proof of LLL (9/9)

- Using the upper bound for the numerator and the lower bound for the denominator, we prove the induction:

$$\Pr[E_k \mid F_S] = \frac{\Pr[E_k \cap F_{S_1} \mid F_{S_2}]}{\Pr[F_{S_1} \mid F_{S_2}]} \leq \frac{p}{1/2} = 2p.$$

- The theorem follows from

$$\begin{aligned} \Pr \left[\bigcap_{i=1}^n \overline{E}_i \right] &= \prod_{i=1}^n \Pr \left[\overline{E}_i \mid \bigcap_{j=1}^{i-1} \overline{E}_j \right] \\ &= \prod_{i=1}^n \left(1 - \Pr \left[E_i \mid \bigcap_{j=1}^{i-1} \overline{E}_j \right] \right) \\ &\geq \prod_{i=1}^n (1 - 2p) > 0. \end{aligned}$$



Outline

- 1 The Lovász Local Lemma
- 2 Applications**
- 3 Explicit Constructions Using the Local Lemma
- 4 Lovász Local Lemma: The General Case



Application: Edge-Disjoint Paths

- Assume that n pairs of users need to communicate using edge-disjoint paths on a given network.
- Each pair $i = 1, 2, \dots, n$ can choose a path from a collection F_i of m paths.
- **Goal:** Apply LLL to show that, if the possible paths do not share too many edges, then there is a way to choose n edge-disjoint paths connecting the n pairs.

Theorem 1

If any path in F_i shares edges with no more than k paths in F_j , where $i \neq j$ and $8nk/m \leq 1$, then there is a way to choose n edge-disjoint paths connecting the n pairs.



Proof of Theorem 1 (1/2)

- Consider the probability space defined by each pair choosing a path independently and uniformly at random from its set of m paths.
- $E_{i,j}$: the event that the paths chosen by pairs i and j **share at least one edge**.
- Since a path in F_i shares edges with no more than k paths in F_j ,

$$p := \Pr[E_{i,j}] \leq \frac{k}{m}.$$



Proof of Theorem 1 (2/2)

- Let d be the degree of the dependency graph. Since event $E_{i,j}$ is independent of all events $E_{i',j'}$ when $i' \notin \{i,j\}$, we have $d < 2n$.
 - Dependency occurs when $E_{i',j}$ or $E_{i,j'}$.
- We have

$$4dp < \frac{8nk}{m} \leq 1,$$

all of the conditions of the LLL are satisfied, proving

$$\Pr\left[\bigcap_{i \neq j} \overline{E}_{i,j}\right] > 0.$$

Hence, there is a choice of paths such that the n paths are edge disjoint.



Application: Satisfiability

Theorem 2

If no variable in a k -SAT formula appears in more than $T = 2^k/4k$ clauses, then the formula has a satisfying assignment.

Proof:

- Consider the probability space defined by the **bad** event E_i :

“the i th clause is **not** satisfied by the random assignment.”

- Since each clause has k literals,

$$\Pr[E_i] = 2^{-k}.$$



Proof of Theorem 2 (2/3)

- The event E_i is **mutually independent** of all of the events related to clauses that **do not share variables** with clause i .
- Because each of the k variables in clause i can appear in no more than $T = 2^k/4k$ clauses, the degree of the dependency graph is bounded by $d \leq kT \leq 2^{k-2}$.
- In this case,

$$4dp \leq 4 \cdot 2^{k-2} 2^{-k} \leq 1.$$



Proof of Theorem 2 (3/3)

- So, we can apply the LLL to conclude that

$$\Pr \left(\bigcap_{i=1}^m \overline{E}_i \right) > 0;$$

hence there is a satisfying assignment for the formula.

Outline

- 1 The Lovász Local Lemma
- 2 Applications
- 3 Explicit Constructions Using the Local Lemma**
- 4 Lovász Local Lemma: The General Case



The issue of LLL on the sampling complexity

- The Lovász Local Lemma proves that a random element in an appropriately defined sample space has a **nonzero probability** of satisfying our requirement.
- However, this probability might be **too small** for an algorithm that is based on simple sampling.
- The number of objects that we need to sample before we find an element that satisfies our requirements might be **exponential** in the problem size.



- In some interesting applications, the Lovász Local Lemma can be used to derive **efficient** construction algorithms.

Common two phases (break the problem into small subproblems)

- 1 a subset of the variables of the problem are assigned random values;
 - the random partial solution fixed in the first phase can be **extended to a full solution** of the problem.
 - the dependency graph H between events defined by the variables deferred to the second phase has, w.h.p., **only small connected components**.
 - 2 the remaining variables are deferred to the second stage.
- Let's see some examples.



Example: A Satisfiability Algorithm (k -SAT)

- **Goal:** a polynomial time algorithm for k -SAT for a constant k .
 - Note: k -SAT is NP-complete for $k \geq 3$.

Input Setting

- Consider a k -SAT formula $\mathcal{F}$, k is an even constant, such that each variable appears in no more than $T = 2^{\alpha k}$ clauses for some constant $\alpha > 0$.
- $x_1, x_2, \dots, x_\ell$: the ℓ variables;
 $C_1, C_2, \dots, C_m$: the m clauses of $\mathcal{F}$.

Dangerous clause

A clause C_i is dangerous if both of the following conditions hold:


- $k/2$ literals of the clause C_i have been fixed;
- C_i is not yet satisfied.

The Two Phases

- **phase I:** Consider the variables $x_1, \dots, x_\ell$ sequentially. If x_i is **not in a dangerous clause**, assign it independently and uniformly at random a value in $\{0, 1\}$.
 - *Surviving clause:* a clause which is not satisfied by variables ($< k/2$) fixed in Phase I.
 - *Deferred variable:* not assigned a value in phase I.
- **phase II:** Exhaustive search to assign values to deferred variables.



From partial assignment to surviving clauses



$$C_1 = (x_1 \vee x_3), C_2 = (\neg x_3 \vee x_4), C_3 = (x_1 \vee \neg x_4),$$

$$C_4 = (x_5 \vee x_6), C_5 = (\neg x_6 \vee x_7), C_6 = (x_2).$$

$$F = C_1 \wedge C_2 \wedge C_3 \wedge C_4 \wedge C_5 \wedge C_6.$$

- x_1, x_2, x_7 : fixed; x_3, x_4, x_5, x_6 : deferred.
- C_6 is satisfied by x_2 so disappears in H' .

The Lemmas

- The partial solution computed in phase I can be extended to a full satisfying assignment of $\mathcal{F}$ (Lemma 1).
- With high probability, the exhaustive search in phase II is completed in time polynomial in m (Lemma 2).

Lemma 1

There is an assignment of values to the deferred variables such that all the surviving clauses are satisfied.

- Existence proof: applying the Lovász Local Lemma.



Proof of Lemma 1 (1/2)

- Let $H = (V, E)$ be a graph on m nodes, where $V = \{1, 2, \dots, m\}$ and let $(i, j) \in E$ if and only if $C_i \cap C_j \neq \emptyset$ (dependency graph).
- Let $H' = (V', E')$ be a subgraph of H such that
 - (a) $i \in V'$ iff C_i is a surviving clause. and
 - (b) $(i, j) \in E'$ iff C_i and C_j share a deferred variable.
- Consider the probability space defined by assigning a random value in $[0, 1]$ independently to each deferred variable.
- Let E_i , for $i = 1, 2, \dots, m$, be the event that surviving clause C_i is NOT satisfied by this assignment (phase I + II).
- Associate E_i with node $i \in V'$. The graph H' is then the dependency graph for this set of events.



Proof of Lemma 1 (2/2)

- A surviving clause has at least $k/2$ deferred variables, so

$$p = \Pr[E_i] \leq 2^{-k/2}.$$


- A variable appears in no more than T clauses; therefore, the degree of the dependency graph is bounded by $d = kT \leq k2^{\alpha k}$.
- For any $k \geq 12$, there is a corresponding suitably small constant $\alpha > 0$ so that

$$4dp = 4k2^{\alpha k}2^{-k/2} \leq 1.$$

- Thus, by the Lovász Local Lemma, there exists an assignment for the deferred variables that (+ the assignment of variables in phase I) satisfies the formula.



Dependency graph and its connected components



Further Observation of the Dependency Graph

- The assignment of values to a subset of the variables in phase I partitions the problem into independent subformulas.
- These subformulas correspond to connected components of the dependency graph H' .
- ★ What if the size of each connected component is **small**?
 $\Rightarrow$ the exhaustive search of all possible assignments can be done efficiently.

Lemma 2

All connected components in H' are of size $O(\log m)$ w.p. $1 - o(1)$.

That is, w.h.p., only small components exist in H' .



Proof of Lemma 2 (1/7)

- The probability that **a given clause survives** is the probability that:
 - either itself is dangerous, or
 - at least one of its direct neighbors is dangerous,

Such a probability is upper bounded by

$$(d + 1)2^{-k/2}, \text{ where } d = kT > 1.$$

- a clause is dangerous with probability $\leq 2^{-k/2}$.
 - union bound.
- We identify a subset $K \subseteq R$ of the vertices in a component R such that the survival of the clauses **represented by the vertices in K are independent events**.



4-Tree of a connected component


4-Tree

A 4-tree S of a connected component R in H is defined as follows:

- ① S is a rooted tree;
- ② any two nodes in S are at distance ≥ 4 in H ;
- ③ there can be an edge in S only between two nodes with distance exactly 4 between them in H ;
- ④ any node of R is either in S or is at distance ≤ 3 from a node in S .

Intuition: A 4-tree covers or dominates a component within radius 3.





◐ vertex in S (4-tree) ◯ vertex in $R \setminus S$

———— edge of S (between nodes at dist. 4 in H)

———— edge of R (edge of H)

Fig.: A 4-tree $S = \{c_0, c_4, c_8\}$ inside a connected component R of H .

Intuitive Idea(s)

- A node u in a 4-tree survives and the event that another node v in a 4-tree survives are **independent**.
 - Any clause that could cause u to survive has distance ≥ 2 from any clause that could cause v to survive.
 - Clauses at distance 2 share no variables, and hence the events that they are dangerous are independent.
- We can take advantage of this independent to conclude that, for any 4-tree S , the probability that the nodes in the 4-tree survive is

$$\leq ((d+1)2^{-k/2})^{|S|}.$$



Proof of Lemma 2 (3/7)

- A **maximal** 4-tree S of a connected component R is the 4-tree with the **largest possible number of vertices**.
- Since the degree of the dependency graph is bounded by d , there are $\leq d + d(d-1) + d(d-1)(d-1) \leq d^3 - 1$ nodes at distance ≤ 3 from any given vertex.
- **Claim:** A maximal 4-tree of R must have $\geq r/d^3$ vertices (note: $r = |V(R)|$).
 - Otherwise, consider the vertices of the maximal 4-tree S and all neighbors within distance ≤ 3 of these vertices, we obtain $< r$ vertices.



Proof of Lemma 2 (4/7): Maximality of S

- Hence, there must be a vertex of distance ≥ 4 from all vertices in S .
- If this vertex has distance exactly 4 from some vertex in S , then it can be added to S and thus S is not maximal ($\Rightarrow \Leftarrow$).
- If its distance > 4 from all vertices in S , consider any path that brings it closer to S ; such a path must eventually pass through a vertex of distance at least 4 from all vertices in S and of distance 4 from some vertex in S ($\Rightarrow \Leftarrow$: maximality of S).



Proof of Lemma 2 (5/7): $1 - o(1)$ constraint

- We show that there is NO 4-tree of H of size r/d^3 that survives with probability $1 - o(1)$.
- Count the number of 4-trees of size $s = r/d^3$ in H . We can choose the root of the 4-tree in m ways.
 - m : number of clauses in the formula $\mathcal{F}$.

The counting argument

- 1 Choose the root of the 4-tree: at most m choices.
- 2 Encode the 4-tree by an Eulerian tour, which has $\leq 2s$ moves.
- 3 Each move corresponds to choosing a length-4 path in H , with $\leq d^4$ choices.



Proof of Lemma 2 (6/7): $1 - o(1)$ constraint

- A tree with root v is **uniquely defined** by an Eulerian tour that starts and ends at v and traverses each edge of the tree twice, once in each direction.
- The counting of 4-trees in H of size s :
 - At most m choices for the root of the 4-tree.
 - An edge of S represents a length-4 path in H .
 - At each vertex in the 4-tree, the Eulerian path can continue in $\leq d^4$ different ways.
 - Thus, the number of 4-trees of size $s = r/d^3$ in H is bounded by $m(d^4)^{2s} = md^{8r/d^3}$.



Proof of Lemma 2 (7/7)

- The probability that the nodes of each such 4-tree survive in H' is

$$\leq ((d+1)2^{-k/2})^s = ((d+1)2^{-k/2})^{r/d^3}.$$

- Hence, by the union bound, the probability that H' has a connected component of size r is bounded by

$$md^{8r/d^3}((d+1)2^{-k/2})^{r/d^3} \leq m2^{(rk/d^3)(8\alpha+2\alpha-1/2)} = o(1)$$

for $r \geq c \log_2 m$ and for a suitably large constant c and a sufficiently small constant $\alpha > 0$.



Solving k -SAT in expected polynomial time for small k

Thus, we have the following theorem.

Theorem 2

Consider a k -SAT formula with m clauses, where k is an even constant and each variable appears in up to $2^{\alpha k}$ clauses for a sufficiently small constant $\alpha > 0$. Then there is an algorithm that finds a satisfying assignment for the formula in expected time that is polynomial in m .

Proof:

- If the Phase I partitions the problem into subformulas involving only $O(k \log m)$ variables (w.p. $1 - o(1)$), then a solution can be found by solving each subformula exhaustively in time polynomial in m .
- Thus, we need only run phase I a constant number of times on average before obtaining a good partition.



Outline

- 1 The Lovász Local Lemma
- 2 Applications
- 3 Explicit Constructions Using the Local Lemma
- 4 Lovász Local Lemma: The General Case**



The General LLL

For completeness we provide the general case of LLL below.

Theorem 2 [Lovász Local Lemma (The General Case); Lovász & Erdős 1975]

Let $E_1, E_2, \dots, E_n$ be a set of events in an arbitrary probability space, and let $G = (V, E)$ be the dependency graph for these events. Assume there exist $x_1, x_2, \dots, x_n \in [0, 1]$ such that, for all $1 \leq i \leq n$,

$$\Pr[E_i] \leq x_i \prod_{(i,j) \in E} (1 - x_j).$$

Then,

$$\Pr \left[\bigcap_{i=1}^n \bar{E}_i \right] \geq \prod_{i=1}^n (1 - x_i).$$



Discussions

